

MetalSynth-Pelvis: Conditioned synthesis of osteosynthesis implants and local metal artifacts in pelvic CT scan volumes using Multi-Window Image-Domain Diffusion

Kiara Alexandra Balcázar Santa Cruz — UTEC · Advisor: Víctor Flores Benites

1. Problem Statement

Automatic pelvic computed tomography (CT) analysis and volumetric segmentation depend on sufficiently annotated data. Metal artifacts can obscure anatomical boundaries near implants, motivating evaluation of segmentation robustness in this region. CTPelvic1K (Liu et al., 2021a) includes a CLINIC-metal subset whose limited bone annotation constrains supervised evaluation. Its name does not establish that every case contains pelvic osteosynthesis, nor that other subsets are free of metal or foreign objects. The local audit will therefore determine implant types, eligible training anatomy and annotated evaluation cases through three-dimensional and multiplanar review. The magnitude of peri-implant segmentation degradation must be measured in this study rather than assumed from the dataset description.

While generative models based on diffusion processes excel at synthesizing visually plausible medical anomalies, they structurally fail when applied to metallic implants. State-of-the-art diffusion-based lesion synthesis models fall into two families, and neither represents the intensity changes that the synthesized object induces outside itself. Mask-bounded inpainting methods condition generation on a lesion mask: CLAIM (Ramzan et al., 2026) recomposes the background from the original image at every denoising step, DiffTumor (Chen et al., 2024) models nothing outside the mask, and LGESynthNet (Jacob et al., 2026) is formulated as inpainting, although its latent compression still perturbs the whole slice incidentally rather than by design. Globally conditioned methods such as DiffBoost (Zhang et al., 2025b) instead synthesize the entire slice from noise, using the mask only as an edge constraint. None of these works states an assumption of tissue non-rigidity; what they share is that no mechanism, and no training signal, accounts for object-induced intensity changes beyond the mask. This is physically incompatible with metal implants, which are rigidly geometric and produce severe attenuation profiles (beam hardening, photon starvation, and high-amplitude streaking) that propagate beyond the implant's boundaries and degrade the signal of the surrounding tissue; reduction networks describe them as present globally in the reconstructed image (Li et al., 2024). De Man et al. (1999) attribute such streaks to beam hardening, scatter, noise and the exponential edge-gradient effect, and show them radiating from the metal along the directions of highest attenuation; their simulation, however, uses an iron rod and dental amalgam in a two-dimensional fan-beam geometry, excludes non-linear partial volume, and reports no spatial extent, so the size of the affected peri-implant band remains unmeasured in the literature.

Physics-based metal artifact simulation provides a reference for appearance generation. This thesis adopts the hybrid simulation protocol of Peters et al. (2025), implemented with CatSim within XCIST (De Man et al., 2007; Wu et al., 2022), as its physical comparison protocol. Their validation compares simulations against a physical phantom scanned on a commercial scanner; it does not cover the hybrid step (clinical images as input, with an empirical frequency-boosting compensation), the vendor-neutral geometry used to generate their data, or pelvic osteosynthesis in three dimensions, and the paper does not state whether patient and metal are projected jointly, which this work verifies in the simulator code before reproduction. Their training data use randomized metal placement, whereas their scoring benchmark uses expert placement. The unresolved placement problem is therefore scalable, anatomically constrained placement, rather than scalable metal insertion itself.

Even if automated placement were achieved, mere anatomical containment is insufficient; a valid simulation must reflect true surgical admissibility. Traditional deterministic surgical planning pipelines (Liu et al., 2025) attempt to optimize a single "perfect" trajectory. Yet real-world clinical data reveal a different reality. Grading iliosacral screws on a four-level cortical-breach scale, Zwingmann et al. (2009) report two distinct postoperative-CT distributions: grade 0 in 69% of navigated placements (grades 1–3: 15%, 8%, 8%) against grade 0 in 40% of conventional placements (37%, 11.5%, 11.5%). That scale reaches iliosacral fixation through Smith et al. (2006), who apply it there, but it originates in the pedicle-screw literature: Gertzbein and Robbins (1990) tabulate canal encroachment from the medial pedicle border in six 2 mm bins rather than four numbered grades, anchor the 2 mm cut to the epidural space, extrapolate the 4 mm cut from a single myelographic CT, and state no justification for the 6 and 8 mm cuts. The four-grade form (0, <2, 2–4, >4 mm) appears literally in Mirza et al. (2003), who declare the thresholds taken from prior studies, measured perforation in cadaveric thoracic spines rather than on CT, and caution that these "safe zone" thresholds do not apply to the thoracic spine and are likely different for different directions of screw perforation. The 2 mm bin width is therefore adopted here as a geometric convention, as with the 10 mm corridor criterion below, and not as a validated clinical threshold; its transfer from the spinal canal to the iliosacral corridor is an explicit assumption of this work. These are two surgical-technique populations that differ significantly, not one aggregate range; this work therefore targets the ordinal distributions themselves rather than a single imperfection rate. The distinction is not merely presentational. In a meta-analysis by the same first author, Zwingmann et al. (2013) pool malposition rates of 2.6% for conventional placement against 0.1% under CT navigation, and the 2009 cohort enters that pool with zero events in both arms, because the included series record malposition only where a revision followed; the authors themselves note that no comparable grading systems are given across studies. Aggregate malposition rates and ordinal cortical-breach grades therefore measure different constructs, and only the latter is a usable target for a placement sampler. Sacral level also warrants separate treatment on anatomical grounds. Using multiplanar CT measurements, Gardner et al. (2010) report mean minimum iliosacral cross-sectional areas at S1 and S2 of 346 and 109.3 mm² in normal sacra, compared with 222 and 220.1 mm² in dysmorphic sacra. The approximately equal areas in the latter group support measuring each level on the individual anatomy rather than assuming a universal corridor-size hierarchy. These measurements concern an uninjured side and ideal central trajectories; they do not describe clinical malposition frequencies. In a retrospective clinical series, van den Bosch et al. (2002) reported postoperative neurologic complaints in 6 of 31 patients with the lower screw in S2, versus 1 of 49 with both screws in S1 ($p = 0.01$). These are patient-level outcomes by screw configuration, not screw-malposition rates by level; the authors also identify a learning-curve bias associated with the change in placement practice. This clinical association supports retaining S1/S2 as a sampler variable, but supplies neither a causal level effect nor an ordinal cortical-breach prior for S2. Cortical breach by level need not follow the same ordering as patient-level neurologic outcome: on postoperative CT, Herman et al. (2017) report breach in 36.5% of S1 screws against 14.8% at S2 ($p = 0.035$), with 32% of screws breaching overall and no screw lying entirely outside bone. The two findings are measured on different quantities and are not in conflict; the latter bounds the range of malpositions the sampler must be able to produce. The ordinal benchmark from Zwingmann et al. (2009) applies to S1: the sacral level is explicit for the conventional arm and not reported for the navigated arm, where it is assumed to be S1, and grades are counted per screw, some patients having received more than one; S2 placements are evaluated descriptively for geometry and breach grades without claiming clinical distribution matching. Its grades were assigned on postoperative CT by one independent radiologist, with no slice thickness, reconstruction interval or inter-observer agreement reported, and the text does not state to which grade a perforation of exactly 2 or 4 mm belongs. Detection itself may depend on image resolution: Tejwani et al. (2014) describe more foraminal penetration on 2.5 mm

than on 5.0 mm sections, although at $P = .3$. Matching against this benchmark is therefore read as agreement between grade distributions measured at an unknown reference resolution, not as voxel-level equivalence, and the grade boundaries used here are a declared convention. The premise that no canonical pose exists is not only inferred from surgical error rates; it is supported anatomically. Measuring the orientation of the bony surfaces that bound the corridor across seventeen cadaveric pelvises, Ziran et al. (2007) report inter-specimen coefficients of variation of 7–25% for most surfaces, rising to 43%, 47–52% and, for the orientation of the superior S1 ala in the sacral frontal plane, 97–140%. A corridor that varies this much between individuals cannot be served by one fixed trajectory. Training a diagnostic network solely on "perfect" deterministic trajectories will inevitably cause it to fail when encountering the true variance of clinical errors. Therefore, a generative sampler must not optimize for perfection, but instead sample from the measured ordinal distribution of clinical malpositions, bounded by pelvic density representations (Arand et al., 2019) and by osseous-corridor viability. For the latter, McLaren et al. (2021) supply the operational element that a qualitative safe-zone description (Ramadanov and Zabler, 2025) does not: a reproducible corridor-diameter procedure over 433 pelvises, yielding transiliosacral angular tolerances of $1.53 \pm 0.57^\circ$ at S1 and $1.02 \pm 0.33^\circ$ at S2. Their ≥ 10 mm corridor criterion is not an empirically validated safety threshold and is not treated as one here: Kaiser et al. (2014), from whom it is taken, describe it as a chosen conservative value and justify it dimensionally, as clearance of 1–2 mm around a 6.3–8 mm screw. It is therefore used as a geometric clearance convention rather than as an inherited clinical standard. Kaiser et al. (2014) additionally supply what a corridor diameter alone does not: a reference frame computable on a CT volume, obtained by reformatting along the sacral axis perpendicular to the superior end plate of S1, with corridor angulation measured coronally against the line joining the iliac crests and axially against the line joining the posterior iliac spines, and a usable-length rule requiring at least 5 mm of clearance to the cortex on either side. The pose distribution sampled within that frame remains an explicit construction of this work.

Recognized Gap. Current attempts to bridge the metal-artifact domain shift are bottlenecked by three mutually exclusive paradigms: physical simulators do not supply scalable anatomically constrained placement, biological diffusion models cannot generate rigid global streaks, and deterministic surgical planners ignore the statistical reality of clinical malposition. To the best of the author's knowledge, there is no diffusion-based method that overcomes these barriers by combining: **(C1)** Non-biological rigid implant geometries, which avoid the threshold-dependent size distortion of fixed-HU metal segmentation (over-coverage in simulated slices, Xie et al., 2024; fragmentation of pelvic screws in the local audit); **(C2)** A surgically constrained placement sampler that accurately emulates the population distribution of clinical malpositions; and **(C3)** The transfer of multi-window HU encoding from artifact *removal*, where it originates, to artifact *synthesis*, coupled with a generation band (B_δ) extended beyond the implant mask so that streaking and beam hardening are modelled in a bounded region outside the metal itself. The multi-window framework is not claimed as novel here: it is established in metal-artifact-reduction work, where published uses apply the windows to per-window reconstruction stages or to the training loss rather than to the encoding of the input (Wang et al., 2025; Li et al., 2024), and Wang et al. (2025) attribute it in turn to earlier multiple-window learning. What is claimed is its use for generation, together with B_δ , for which no published precedent quantifying a *generation band outside the implant mask* was found. Peri-implant artifact extent itself has been measured: Radzi et al. (2014) report CT artifact distances of "2.0, 2.6, 1.6 and 2.0 mm" across titanium, stainless-steel and cannulated screws, and Cassanego et al. (2026) report 3.1–4.2 mm by the same style of measurement. Neither supplies a calibration for B_δ . Both measure the "perpendicular distance from the central screw axis to the boundary of the artifact", that is from the screw *axis* rather than from its surface, to the edge of a thresholded surface whose threshold is not published, on 3.5–4.0 mm screws in single-specimen ankle and dental settings; they do not report the quantity in Hounsfield units, and their two ranges are not mutually comparable for lack of a common reference point. The band width used here is therefore a declared construction, not an inherited measurement. This is the gap this work aims to fill.

Filling it rests on one assumption that the present literature neither grants nor refutes: that the *appearance* of a metal artifact can be generated in the image domain, without access to projection data. The mechanisms involved are non-linear and are manipulated in the sinogram (De Man et al., 1999), and dual-domain reduction networks report a penalty for operating on images alone: Lin et al. (2019) obtain 31.45 dB PSNR from their image-domain variant against 33.51 dB for the full dual-domain network. That image-domain variant is fed a reconstruction of an already-interpolated sinogram, but the penalty does not rest on this detail: Li et al. (2024), whose image-only branch starts from the metal-corrupted reconstruction, also report it below their multi-domain network and state that working in a single domain cannot effectively recover the true tissues. One qualification keeps these results from settling the question: both address artifact *removal*, which Lin et al. (2019) describe as ill-posed because information has been destroyed, whereas synthesis carries no such recovery requirement. This work therefore treats image-domain generability not as a premise but as part of what its distributional physical-coherence evaluation puts to the test, with the physics-based protocol of Peters et al. (2025) supplying the mechanism-faithful reference against which the generated appearance is judged.

2. Research Question

What improvement does a surgical admissibility-constrained (bone density and cortical containment) 3D placement sampler, a multi-window image-domain diffusion model conditioned on the rigid implant geometries and an extended generation band (B_δ), have in terms of distributional physical-coherence of the synthesized pelvic osteosynthesis implants and their local artifacts, when compared with both naive copy-paste insertion and the adopted physics-based simulation protocol of Peters et al. (2025)? Downstream segmentation impact is deliberately excluded from this study and is stated as out of scope below.

3. Hypothesis

The synthetic generation of medical anomalies and metallic implants using a rigid geometry and extended spatial boundaries, HU-conditioned diffusion model and whose placement is surgically enforced via a density-based placement sampler is more physically-coherent than unconstrained generative baseline methods and on the same footing as the adopted explicit physics-based simulation protocol (Peters et al., 2025). Placement coherence is quantified by Surgical Admissibility of Placement (SAP), the single metric introduced by this work, which emulates the measured ordinal distribution of real-world implant placement. Appearance coherence is quantified with the metrics of the adopted protocol under their published names, not with newly coined ones.

4. Objectives

General Objective. Design, implement and evaluate a physics-and-anatomy-aware image-domain diffusion pipeline with rigid surgical constraints and extended HU-based conditioning, capable of generating physically coherent metallic implants intended, in future work, to improve the robustness of downstream segmentation networks.

1. **Validation of a multi-window representation (Compulsory Go/No-Go):** Ensure HU reconstruction error (HU to multi-window to VAE to HU again) within the bone tissue, delimited by the same 150 HU bone threshold used by the adopted evaluation protocol (Peters et al., 2025; Haneda et al., 2025), below a 25 HU threshold (Mean Absolute Error (MAE)). No published pass/fail threshold on CT-number accuracy was found, so this criterion is not taken from a source: it is anchored on the error level that current metal-artifact methods reach on the same scale. The NMAR algorithm that calibrates the community benchmark reports an RMSE of 20.2 HU, a sinogram-domain diffusion model evaluated in image space 12.3 HU (Karageorgos et al., 2024), and a latent diffusion model 12.74 HU (Yun et al., 2026). Because those values are RMSE while this criterion is MAE, and RMSE is never below MAE for the same residuals, the comparison fixes the order of magnitude rather than an equivalence; this is declared rather than resolved. A concurrent preprint reports that frozen video autoencoders transfer to CT without medical fine-tuning (Chen et al., 2026); it reports PSNR, SSIM and MSE without declared units and no HU error, evaluates no bone or metal region, uses no pelvic or metal-bearing volumes, and declares intensity clipping to $[-1000, 1000]$ HU for its generation experiments, so it does not address the dense-bone and metal range this criterion targets. In volumes with metal, the same round-trip error is additionally reported inside the band B_δ around the metal, descriptively and outside the Go/No-Go: that band holds the low-amplitude streaks the synthesizer must render, and an error restricted to bone does not show whether the autoencoder preserves them, since the first stage of latent diffusion is designed to remove high-frequency detail (Rombach et al., 2022). The gate rule is fixed before any run. It is evaluated on 34 held-out patients (patient-level split stratified by dataset and metal presence) for the pretrained autoencoder and for a decoder-adapted variant with the encoder frozen, each under three three-channel encodings: the windows reported by Wang et al. (2025) $[-1000, 2000]$, $[-320, 480]$ and $[-160, 240]$ HU, the same with the wide-window ceiling raised to 20000 HU, and the wide window replaced by an arcsinh compression over $[-1000, 20000]$ HU. HU are read back without ground truth, from the narrowest non-saturated channel. A combination passes when the mean over test patients of the per-patient bone MAE is below 25 HU, and the gate passes if any of the six combinations does. This multiplicity is not corrected; the patient-level bootstrap 95% confidence interval is reported with each verdict, and a pass whose upper bound reaches 25 HU is reported as marginal. If several combinations pass, the one carried to Objective 3 follows an order fixed a priori rather than their errors: arcsinh encoding, then the 20000 HU ceiling, then the published windows, whose 2000 HU ceiling clips metal; within an encoding, the pretrained decoder precedes the adapted one. If none passes, the latent route is abandoned and the No-Go result is reported as the outcome. *Measured outcome.* The gate was run under this rule. On the 34 held-out patients the pretrained autoencoder reached a mean per-patient bone MAE of 160.13–209.39 HU and the decoder-adapted variant 61.72–77.03 HU; the best of the six combinations, arcsinh encoding with an adapted decoder, gave 61.72 HU (bootstrap 95% CI [55.12, 68.99]), with all 34 patients above the threshold in every combination. Identity controls reproduced the window-only reference (165/165) and the pretrained round-trip reproduced an earlier full-cohort run (330/330); encoder hashes confirm it stayed frozen and the validation curves had flattened, so the result is not an artefact of insufficient training. Inside B_δ the same adapted decoder left 69.89 HU, so restricting the latent to the generation band would not recover the loss either. The outcome is therefore **No-Go** and is reported as a result: pretrained natural-image latents do not transport dense-bone CT numbers at this tolerance. A CT-trained latent was examined next, under the same rule and on the same held-out patients, and was ruled out before running it: the released MAISI bundle (Guo et al., 2025) maps its output to $[-1000, 1000]$ HU in its own sampling script, a range the article itself does not state, so dense cortical bone above that ceiling and metal above 2500 HU fall outside the representation by construction. Clipping the held-out volumes to that range, with no model involved, already yields a mean per-patient bone MAE of 42.24 HU, above the 25 HU criterion in 11 of the 34 patients; since this is a lower bound that no choice of weights can improve, the evaluation was not run. The finding generalizes the outcome rather than being incidental to one architecture: the pretrained latents available here, including one trained on CT, are defined over intensity ranges that exclude dense bone and metal. Objective 3 responds by synthesizing in the image domain; that redesign was decided after seeing this result and is declared as such, while the gate rule itself was fixed beforehand and is not modified. This will serve as a guarantee that the intensity gradients that drive downstream applications are preserved.
2. **Design of surgically constrained placement sampler:** propose position-orientation 3D pose of a parametric rigid screw model constrained by bone density read from the CT numbers of each volume, motivated by the systematic heterogeneity of pelvic bone mass (e.g., low grey values in the sacral ala) shown by a population-mean statistical model (Arand et al., 2019), cortical boundaries with a 5 mm clearance rule, and osseous-corridor viability (McLaren et al., 2021), expressed in the CT-computable reference frame of Kaiser et al. (2014) (sacral-axis reformat; coronal and axial corridor angulation against iliac-crest and posterior-iliac-spine landmarks). Corridor geometry is measured on anatomical masks from TotalSegmentator 2.18.0 (Wasserthal et al., 2023), which is built on nnU-Net (Isensee et al., 2021) (task `total`, default crop; robust crop as sensitivity analysis), quality-controlled per case and cleaned of connected components below 0.1% of their structure, rather than on an HU threshold. The cited publication evaluates a 104-structure model and reports no accuracy for the S1 label used here, so its reported accuracy is not assumed for the version and classes used in this work. The sampler must yield a distribution of placements emulating the two measured ordinal malposition distributions of Zwingmann et al. (2009), conditioned on surgical technique within S1, rather than a single artificial "optimized" trajectory. Sacral level remains a geometric conditioning variable, anatomically motivated by the segment-specific corridor measurements of Gardner et al. (2010); S2 results are reported descriptively, because the clinical comparison of van den Bosch et al. (2002) concerns neurologic complaints by screw configuration and does not provide an ordinal S2 placement distribution. Sacral-morphology phenotypes are deliberately not used as a stratifying variable: their reported effect on safe-zone size is not consistent across studies, and the sampler instead measures corridor geometry directly on each volume.
3. **Design of a conditioned synthesizer:** 2.5D conditional diffusion in the image domain, formulated as inpainting of the generation region $G = M \cup B_\delta$, where M is the rigid implant mask. The model receives the patch with G removed, the implant mask and the band mask, and generates CT numbers inside G only; everything outside G is copied from the source volume, so preservation outside the band holds by construction rather than by measurement. This replaces the latent formulation and, with it, the ControlNet-guided Stable Diffusion 1.5 backbone: ControlNet presupposes a frozen pretrained base and its latent space (Zhang et al., 2023), Objective 1 measured that this latent does not preserve bone CT numbers at the required tolerance, and no pretrained pixel-space CT base is adopted here. The multi-window encoding is kept as the input and output channel representation, now with no compression stage between them. A crucial detail in the formulation is that the synthesis mask must be extended with a spatial band B_δ ($\sim 12\text{mm}$) over the implant boundary, to allow a more explicit manifestation of the streaking artifacts and beam hardening that occur outside of the implant mask in the generated image. The band width is a design parameter: it bounds where artifact is generated and deliberately truncates far-field streaks, whose spatial extent remains unmeasured in the reviewed literature. No published value calibrates it, but its direction is supported: in a metal-mask sensitivity study, Karageorgos et al. (2024) report that dilating the metal trace leaves reconstruction error essentially unchanged (RMSE 11.45 and 11.63 HU at dilation factors of 1.4 and 1.8), whereas shrinking it makes error collapse (54.82 and 57.69 HU at factors of 0.7 and 0.5), against 7.57 HU for the true trace. Under-covering the extent of metal is therefore far more damaging than over-covering it, which is the

argument for generating beyond the implant mask rather than inside it. That experiment varies a mask in sinogram-domain artifact reduction, not a generation band in the image domain, so it motivates the direction of B_δ without fixing its width; the 12 mm value remains a declared construction. Training pairs come from the metal-bearing volumes of the training patients, where the real implant supplies both the mask and the appearance target, under the same patient-level isolation as the rest of the work; at synthesis time the mask is the parametric screw of Objective 2 placed in a sampled pose. The training unit is the individual connected metal component rather than the axial slice, so that a training region has the same structure as a synthesis region, which contains a single implant. Inclusion is governed by a criterion fixed before training: the component must lie inside the patient body, the case must carry declared orthopaedic material, and the patch must contain metal of the target component, together with a declared proportion of band-only patches. Component size and shape are reported, not filtered, because the model learns a mask-to-artifact relation that does not depend on implant type. Three domain shifts are declared rather than assumed away: training masks are obtained by thresholding real metal, whose shaft the local audit measured below the published calibres, while the synthesis mask is a smooth parametric cylinder; the context outside G carries far-field streaks in metal-bearing training volumes but not in the metal-free volumes used for synthesis; and band-only patches, which contain generation region but no metal, are less frequent in the extracted training set than the corridor geometry predicts for synthesis. The first and third are quantified and reported; the second is evaluated through the continuity of CT numbers across the band boundary.

4. **Formalization of the placement metric and adoption of an inherited appearance protocol:** Formulation of Surgical Admissibility of Placement (SAP) as the only metric introduced by this work, evaluated with Wasserstein-1 distances against the measured ordinal distributions. Appearance is evaluated with the metrics of the adopted protocol under their published names — bone integrity, metal integrity and streak amplitude (Peters et al., 2025) — and the previously proposed BFC and ISC are withdrawn, so that inherited measures are not renamed as contributions.

Explicitly out of scope. Downstream segmentation evaluation (Dice, HD95) is excluded from this study, and constraint-by-constraint ablations are deferred to future work. Two auditable conditions drive the exclusion rather than time alone: within the locally available cohort only a minority of the metal-bearing volumes carry verified bone annotation, and only part of the published collection is held locally, so a downstream result would rest on a cohort too small and too partially annotated to support the claim. The synthesis arms and the placement sampler are therefore evaluated on physical coherence and surgical admissibility alone; downstream utility is left as future work.

5. Expected Results

Goal	Metric(s)	Expected Outcome
Surgical Admissibility (Only novel metric)	SAP: four-level Cortical Breach Grade (Smith et al., 2006), with boundaries fixed by convention as grade 0 for no perforation, 1 for (0, 2), 2 for [2, 4] and 3 for (4, ∞) mm; density-zone fraction; corridor viability (McLaren et al., 2021) in the reference frame of Kaiser et al. (2014).	Distribution matching: Wasserstein-1 distance shows generated placements approaching the two measured ordinal distributions of Zwingmann et al. (2009), conditioned on technique within S1. S2 geometry and breach grades are reported descriptively, without a clinical ordinal target.
Appearance Coherence (Inherited protocol, published names)	Bone integrity; metal integrity; streak amplitude (Peters et al., 2025). Inversion of each metric for synthesis is documented, and the 0–4 scale is not carried over.	Streak amplitude is the single primary endpoint, declared before the test set is examined; the remaining metrics are descriptive. Superiority over naive copy-paste insertion is assessed with a one-sided paired Wilcoxon signed-rank test, and comparability with the physics-based protocol with a two one-sided test of equivalence, so that equivalence is claimed only when the paired difference lies within a margin fixed in advance, never inferred from a non-significant difference. That margin is set from the method’s own variability under conditions that should not change the result, measured on validation patients only. With the available sample sizes a non-conclusive equivalence result is a possible and reportable outcome.
Representation Viability (Compulsory gate, executed)	MAE in HU (Round-trip; bone, and B_δ descriptively)	Result: No-Go. Best of the six pre-specified combinations: 61.72 HU (95% CI [55.12, 68.99]), with 34/34 held-out patients above the threshold. The criterion was a mean per-patient bone error strictly below 25 HU, a level anchored on the CT-number accuracy reported for NMAR (20.2 HU) and for learned metal-artifact methods (12.3–12.74 HU) (Karageorgos et al., 2024; Yun et al., 2026), not on a published pass/fail criterion. Log-compressed metal ranges preserve the structural topology.
Preservation outside B_δ (No-damage check)	RMSE, SSIM against the original CT, restricted to the complement of B_δ .	No measurable change where none is intended; deviation inside B_δ is the intended signal and is not scored as error.

6. Datasets

Primary Dataset. CTPelvic1K (Liu et al., 2021a) will be used as the source of patient anatomy and real artifact observations. The locally available dataset6 and dataset7 volumes are inventoried separately from the full published collection; the CLINIC-metal label alone does not establish implant type or annotation availability. Acquisition and reconstruction settings, including any vendor metal artifact reduction or monoenergetic reconstruction, are not reported in the dataset publication; since both change artifact appearance (Selles et al., 2024), real artifact observations are treated as an appearance reference of unknown reconstruction, not as a physical ground truth. *Strict Isolation Rule.* No patient may contribute to both training and evaluation. The implant geometries consumed by the proposed sampler and by the physical simulation arm are parametric and are not derived from any evaluation volume. Duplicate-volume screening and patient grouping are required to enforce this separation.

Implant geometry source. Implants are modelled as parametric rigid screws rather than extracted from the metal-bearing volumes. Two geometries are kept apart, because they answer different questions. The *corridor envelope* used for placement viability is the nominal calibre reported for iliosacral fixation, 6.5–8.0 mm (Gardner et al., 2010) and 6.3–8 mm (Kaiser et al., 2014). The *metal mask* used for synthesis is the body of the implant, which is thinner: the manufacturer’s technique guide for this system prints a 4.8 mm shaft and states that core and shaft diameters coincide (Syn, 2003), a biomechanical study reports 4.9 mm shaft and 4.7 mm minor diameter under the thread for the same

nominal calibre (Gardner et al., 2015), and a finite-element study of the same screw uses 7.3 mm only for the distal thread, with a 2.5 mm pitch over 16 mm, and 4.8 mm for the shaft (Zhu and Liao, 2023). The nominal figure is therefore the thread diameter, not the body, and a uniform 6.5–8.0 mm cylinder would overstate the metal cross-section along the corridor by more than a factor of two. The head is modelled from the only published measurement found, 8.0 mm in diameter and 4.5 mm high, seated without countersinking (Sayres et al., 2015). Screw length is not taken from a fixed published range: it is bounded on each volume by the measured corridor under the usable-length rule of Kaiser et al. (2014). Three parameters remain without a source in the manufacturer’s own documentation and are declared rather than assumed: washer thickness, taken as 1.5 mm from a different manufacturer’s catalogue (Dou, 2021) against an outer diameter of 13.0 mm that three sources agree on; the cannulation diameter, which is left as a free parameter of the model and evaluated solid against hollow, since the 2.9 mm value circulating in distributor listings appears in the manufacturer guide only as the cannulation of drills and screwdrivers; and the alloy, since the same system is supplied in 316L stainless steel and in Ti-6Al-7Nb (Syn, 2003) while a binary mask cannot encode which one, so the renderer learns the appearance of a mixture of materials. The real implants of the metal-bearing subset do not enter the geometry bank; they serve as a descriptive contrast for placements, as an appearance reference for synthesis, and as a test–retest control where the same implant was imaged twice. The local audit motivates this choice: screws extracted from these volumes with a fixed HU threshold appear fragmented, with shaft diameters below the published calibres, and the extracted geometry changes between acquisitions of the same implant.

Physical comparison protocol. The hybrid simulation and evaluation protocol of Peters et al. (2025), also associated with the AAPM CT-MAR Challenge (Haneda et al., 2025), is adopted instead of an independently validated reimplement of XCIST/CatSim. Challenge data remain excluded from model training, and the scoring benchmark is reserved for final evaluation. Reproduction will document the code revision, configuration, reconstruction settings and access conditions. The physical arm is run on a reduced evaluation subset of patients. Both synthesis arms will use the same source anatomy and implant poses; a no-metal simulation/reconstruction control will quantify changes introduced by the physical arm alone. The published protocol is two-dimensional and evaluates MAR: extending it to pelvic osteosynthesis synthesis requires an explicit adaptation and does not inherit its validation automatically. RMSE and SSIM against the original CT will assess preservation outside the generation band, where changes are not intended; artifact realism within the affected region requires separate evaluation. No unverified numerical acceptance thresholds or MAR score calibration are adopted here.

Field limitation: corridor geometry has almost only been characterized in intact pelves. This is not a limitation of the present cohort but of the available literature. With one exception, every source that publishes sacral corridor geometry measured it in pelves without implants and without fracture: Kaiser et al. (2014) excluded scans with implants obscuring the lumbosacral junction, McLaren et al. (2021) used uninjured scans, and Ziran et al. (2007) worked on cadaveric specimens and warn that fracture, displacement and soft-tissue injury would alter what is visible. The exception, Reilly et al. (2003), osteotomised a zone II sacral fracture in six cadaveric pelves and found that 5–20 mm of cranial displacement reduced the cross-sectional area available to S1 iliosacral screws by 36–90%, so corridor limits obtained on intact pelves do not transfer to fractured or malreduced ones; this supports measuring the corridor on each volume rather than importing a published limit. It also exposes an uncontrolled difference between benchmark and host: the clinical ordinal distributions of Zwingmann et al. (2009) come from fractured pelves (Tile types B and C), whereas corridors here are measured on pelves without osteosynthesis, so any residual malreduction in the clinical cohort narrows its corridors in a way the synthetic hosts do not reproduce. The reference frame adopted here rests on landmarks at the lumbosacral junction, which is precisely where streaking concentrates, so whether those landmarks remain locatable in metal-bearing volumes is not answered by any published source. It was therefore quantified on the local cohort. Landmarks were located by a heuristic on each volume, and a landmark region was counted as contaminated when its dark-streak or metal fraction exceeded the maximum observed in 69 metal-free patients of the same cohort. Among 65 patients with orthopaedic metal, all five landmarks were found within the field of view in 57. After a clinician reviewer (one surgeon, otorhinolaryngology), blinded to the automatic audit, confirmed the S1 level on every patient, the frame was computable in 48 patients and free of contamination at all five landmark regions in 29. Losses were driven more by field-of-view truncation of the iliac crests (7 patients) and by landmark localisation than by artifact contamination; in the 17 patients with a screening-detected iliosacral screw, the S1 level was correct in all and its region was contaminated in 3. The frame is therefore computed on metal-bearing volumes where it is computable; the handling of field-of-view-truncated cases is assessed separately, and the heuristic localisation of the iliac crests and posterior iliac spines has not been clinically reviewed.

Corridor measurement on the local cohort. Among 103 patients without osteosynthesis, 91 had a located S1 and 72 passed level quality control. Level discordance excluded 11 of 34 patients with non-orthopaedic objects and 7 of 57 clean patients, and one further patient was excluded because S1 reached the field-of-view border; this imbalance is declared. The maximum transsacral corridor diameter had a median of 9.5 mm (IQR 7.4–11.7), and the 10 mm convention was met in 29 of 72 patients with the default crop and in 27 with the robust crop; the crop is therefore reported as a measurement uncertainty that reaches the binary viability outcome. Component cleaning changed neither the transsacral corridor diameter nor its 10 mm viability in this cohort at any tested fraction up to 5%; the fraction, like the occupancy rule below, was fixed after inspecting these results. Along the axis of the widest corridor, the median fraction of voxels at or below 150 HU inside the sacral masks was 0.42, so an HU-based bone definition would exclude much of the corridor. In metal-bearing volumes, where the segmentation labels the existing implant as bone, treating the implant as occupied space (local half-maximum threshold per object) lowered 10 mm viability from 53.8% to 40.4% in 52 patients with the default crop, and a fixed 2500 HU threshold, which thins the screws and therefore bounds the occupied corridor from above, gave the same proportion; these volumes are reported as a separate contrast and are not pooled with the corridor cohort.

Local cohort audit. Training eligibility requires manual review of three-dimensional views and multiplanar CT slices confirming absence of metal and foreign objects. HU screening only prioritizes this review, including cases from dataset6. Confirmed metal cases are candidates for testing, with pelvic osteosynthesis distinguished from other implants and external objects. Exact volume duplicates and repeated patients must be resolved before splitting. Final cohort sizes remain pending this audit.

Why downstream evaluation is excluded. The audit itself supplies the reason. Only a minority of the metal-bearing volumes carry verified bone annotation, and only a fraction of the published collection is available locally, so a Dice or HD95 result would rest on a cohort too small and too partially annotated to support a robustness claim. The metal-bearing subset used here is nonetheless the de facto clinical evaluation set for pelvic metal-artifact work, having been used as such in recent published work (Wang et al., 2025). Where a clean reference image does not exist for these volumes, appearance realism is assessed against real artifact observations rather than against an unavailable ground truth.

7. References

- E. Abadi, B. Harrawood, S. Sharma, A. Kapadia, W. P. Segars, and E. Samei. DukeSim: A Realistic, Rapid, and Scanner-Specific Simulation Framework in Computed Tomography. *IEEE Transactions on Medical Imaging*, 38(6):1457–1465, 2019. doi: 10.1109/TMI.2018.2886530.
- 6.5 mm/7.3 mm Cannulated Screw System. *Surgical Technique*. Acumed LLC, Hillsboro, OR, USA, 2020. Document SPF10-05-E, effective 2020/09.
- J. An, D. Wang, P. Guo, J. Luo, and A. Schwing. On Inductive Biases That Enable Generalization in Diffusion Transformers. In D. Belgrave, C. Zhang, H. Lin, R. Pascanu, P. Koniusz, M. Ghassemi, and N. Chen, editors, *Advances in Neural Information Processing Systems*, volume 38, Main Conference, pages 31738–31757. Curran Associates, Inc., 2025. doi: 10.52202/085713-1066.
- C. Arand, D. Wagner, R. G. Richards, H. Noser, L. Kamer, T. Sawaguchi, and P. M. Rommens. 3D Statistical Model of the Pelvic Ring—A CT-Based Statistical Evaluation of Anatomical Variation. *Journal of Anatomy*, 234(3):376–383, 2019. doi: 10.1111/joa.12928.
- T. Berk, I. Zderic, P. Schwarzenberg, T. Pastor, S. Halvachizadeh, G. Richards, B. Gueorguiev, and H.-C. Pape. Is a Washer a Mandatory Component in Young Trauma Patients with S1-S2 Iliosacral Screw Fixation of Posterior Pelvis Ring Injuries? A Biomechanical Study. *Medicina*, 59(8):1379, 2023. doi: 10.3390/medicina59081379.
- G. R. Cassanego, S. C. Rahal, T. S. G. Santos, J. P. d. Silva, D. B. R. d. Santos, M. J. Mamprim, M. J. Carbonari, and R. C. Siqueira. Evaluation of Artifacts Produced by Cannulated and Solid-Core Screws Using Magnetic Resonance Imaging and Computed Tomography. *The Veterinary Journal*, 317:106691, 2026. doi: 10.1016/j.tvjl.2026.106691.
- Q. Chen, X. Chen, H. Song, Z. Xiong, A. Yuille, C. Wei, and Z. Zhou. Towards Generalizable Tumor Synthesis. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11147–11158, 2024. doi: 10.1109/CVPR52733.2024.01060.
- Q. Chen, S. Ding, Y. Gu, N. Liu, J. Bian, A. Yuille, Z. Zhou, and J. Fu. Foundation VAEs for 3D CT Reconstruction, Augmentation, and Generation, 2026.
- D.-I. Choi, S. Yun, S. Hyun, and S. Cho. Metal Artifact Reduction Algorithm with Conditional Latent Diffusion Model for Dental Cone-Beam CT. *Journal of Applied Clinical Medical Physics*, 26(11): e70317, 2025. doi: 10.1002/acm2.70317.
- B. De Man, J. Nuyts, P. Dupont, G. Marchal, and P. Suetens. Metal Streak Artifacts in X-Ray Computed Tomography: A Simulation Study. *IEEE Transactions on Nuclear Science*, 46(3):691–696, 1999. doi: 10.1109/23.775600.
- B. De Man, S. Basu, N. Chandra, B. Dunham, P. Edic, M. Iatrou, S. McOlash, P. Sainath, C. Shaughnessy, B. Tower, and E. Williams. CatSim: A New Computer Assisted Tomography Simulation Environment. In J. Hsieh and M. J. Flynn, editors, *Medical Imaging 2007: Physics of Medical Imaging*, volume 6510, page 65102G. SPIE, 2007. doi: 10.1117/12.710713.
- P. Dhariwal and A. Nichol. Diffusion Models Beat GANs on Image Synthesis. In *Proceedings of the 35th International Conference on Neural Information Processing Systems, NIPS '21*, Red Hook, NY, USA, 2021. Curran Associates Inc. ISBN 9781713845393.
- Z. Dorjsenbe, H.-K. Pao, and F. Xiao. Polyp-DDPM: Diffusion-Based Semantic Polyp Synthesis for Enhanced Segmentation. In *2024 46th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 1–7, 2024. doi: 10.1109/EMBC53108.2024.10782077. *Trauma Catalogue*. Double Medical, 2021. Edition 2021.
- N. A. Ebraheim, R. J. Coombs, M. J. Hoefflinger, and W. T. Jackson. A Pitfall of Radiologic Evaluation of Sacroiliac Joint Screw Positioning. *Orthopedics*, 16(5):616–618, 1993. doi: 10.3928/0147-7447-19930501-16.
- N. A. Ebraheim, R. Xu, A. Biyani, and M. C. Nadaud. Morphologic Considerations of the First Sacral Pedicle for Iliosacral Screw Placement. *Spine*, 22(8):841–846, 1997. doi: 10.1097/00007632-199704150-00002.
- N. A. Ebraheim, S. P. Haman, R. Xu, S. Stanescu, and R. A. Yeasting. The Lumbosacral Nerves in Relation to Dorsal S1 Screw Placement and Their Locations on Plain Radiographs. *Orthopedics*, 23(3):245–247, 2000. doi: 10.3928/0147-7447-20000301-20.
- Y. Fan, J. Pack, and B. De Man. A Virtual Imaging Trial Framework to Study Cardiac CT Blooming Artifacts. In *7th International Conference on Image Formation in X-Ray Computed Tomography*, volume 12304 of *Society of Photo-Optical Instrumentation Engineers (SPIE) Conference Series*, page 1230415, 2022. doi: 10.1117/12.2646407.
- A. W. Gardner, M. Z. Toh, K. S. A. Yew, D. T. T. Lie, and S. M. Chou. Cannulated Versus Non-Cannulated Cancellous Screw Fixation for Femoral Neck Fractures: A Synthetic Bone Biomechanical Study. *Journal of Orthopaedic Surgery*, 23(1):41–46, 2015. doi: 10.1177/230949901502300110.
- M. J. Gardner and M. L. C. Routt, Jr. Transiliac-Transsacral Screws for Posterior Pelvic Stabilization. *Journal of Orthopaedic Trauma*, 25(6):378–384, 2011. doi: 10.1097/BOT.0b013e3181e47fad.
- M. J. Gardner, S. Morshed, S. E. Nork, W. M. Ricci, and M. L. Chip Routt, Jr. Quantification of the Upper and Second Sacral Segment Safe Zones in Normal and Dysmorphic Sacra. *Journal of Orthopaedic Trauma*, 24(10):622–629, 2010. doi: 10.1097/BOT.0b013e3181cf0404.
- S. D. Gertzbein and S. E. Robbins. Accuracy of Pedicular Screw Placement In Vivo. *Spine*, 15(1):11–14, 1990. doi: 10.1097/00007632-199001000-00004.
- H. Gottschling, M. Schröder, N. Reimers, F. Fischer, A. Homeier, and R. Burgkart. A System for Performing Automated Measurements on Large Bone Databases. In O. Dössel and W. C. Schlegel, editors, *World Congress on Medical Physics and Biomedical Engineering, September 7 - 12, 2009, Munich, Germany*, pages 910–913, Berlin, Heidelberg, 2010. Springer Berlin Heidelberg. ISBN 978-3-642-03882-2.
- F. Gras, H. Gottschling, M. Schröder, I. Marintschev, G. O. Hofmann, and R. Burgkart. Transsacral Osseous Corridor Anatomy Is More Amenable To Screw Insertion In Males: A Biomorphometric Analysis of 280 Pelves. *Clinical Orthopaedics and Related Research*, 474(12):2304–2311, 2016. doi: 10.1007/s11999-016-4954-5.
- D. R. Griffin, A. J. Starr, C. M. Reinert, A. L. Jones, and S. Whitlock. Vertically Unstable Pelvic Fractures Fixed with Percutaneous Iliosacral Screws: Does Posterior Injury Pattern Predict Fixation Failure? *Journal of Orthopaedic Trauma*, 20(1 Suppl):S30–S36, 2006.
- P. Guo, C. Zhao, D. Yang, Z. Xu, V. Nath, Y. Tang, B. Simon, M. Belue, S. Harmon, B. Turkbey, and D. Xu. MAISI: Medical AI for Synthetic Imaging. In *2025 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 4430–4441, 2025. doi: 10.1109/WACV61041.2025.00435.
- E. Haneda, N. Peters, J. Zhang, G. Karageorgos, W. Xia, H. Paganetti, G. Wang, Y. Guo, J. Ma, H. S. Park, K. Jeon, F. Fan, M. Thies, and B. De Man. AAPM CT Metal Artifact Reduction Grand Challenge. *Medical Physics*, 52(10):e70050, 2025. doi: 10.1002/mp.70050.
- E. A. Hasenboehler, P. F. Stabel, A. Williams, W. R. Smith, J. T. Newman, D. L. Symonds, and S. J. Morgan. Prevalence of Sacral Dysmorphism in a Prospective Trauma Population: Implications for a “Safe” Surgical Corridor for Sacro-Iliac Screw Placement. *Patient Safety in Surgery*, 5(1):8, 2011. doi: 10.1186/1754-9493-5-8.
- A. Herman, E. Keener, C. Dubose, and J. A. Lowe. Simple Mathematical Model of Sacroiliac Screws Safe-Zone-Easy to Implement by Pelvic Inlet and Outlet Views. *Journal of Orthopaedic Research*, 35(7):1478–1484, 2017. doi: 10.1002/jor.23396.
- A. F. Hinsche, P. V. Giannoudis, and R. M. Smith. Fluoroscopy-Based Multiplanar Image Guidance for Insertion of Sacroiliac Screws. *Clinical Orthopaedics and Related Research*, (395):135–144, 2002. doi: 10.1097/00003086-200202000-00014.
- J. Ho, A. Jain, and P. Abbeel. Denoising Diffusion Probabilistic Models. In *Proceedings of the 34th International Conference on Neural Information Processing Systems, NIPS '20*, Red Hook, NY, USA, 2020. Curran Associates Inc. ISBN 9781713829546.
- Q. Hu, Y. Chen, J. Xiao, S. Sun, J. Chen, A. Yuille, and Z. Zhou. Label-Free Liver Tumor Segmentation. In *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7422–7432, 2023. doi: 10.1109/CVPR52729.2023.00717.
- F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein. nnU-Net: A Self-Configuring Method for Deep Learning-Based Biomedical Image Segmentation. *Medical Methods*, 18(2):203–211, 2021. doi: 10.1038/s41592-020-01008-z.
- A. J. Jacob, P. Sharma, and D. Rueckert. LGESynthNet: Controlled Scar Synthesis for Improved Scar Segmentation in Cardiac LGE-MRI Imaging. In O. Camara, E. Puyol Antón, M. Sermesant, C. Mauge, M. Varela, Y. Ma, R. Paulsen, C. Wang, Q. Tao, and A. Young, editors, *Statistical Atlases and Computational Models of the Heart. Regular and CMRxRecon Challenge Papers*, pages 34–44, Cham, 2026. Springer Nature Switzerland. ISBN 978-3-032-17734-6. doi: 10.1007/978-3-032-17734-6_4.
- S. P. Kaiser, M. J. Gardner, J. Liu, M. L. C. Routt, Jr., and S. Morshed. Anatomic Determinants of Sacral Dysmorphism and Implications for Safe Iliosacral Screw Placement. *The Journal of Bone and Joint Surgery. American Volume*, 96(14):e120, 2014. doi: 10.2106/JBJS.M.00895.
- G. M. Karageorgos, J. Zhang, N. Peters, W. Xia, C. Niu, H. Paganetti, G. Wang, and B. De Man. A Denoising Diffusion Probabilistic Model for Metal Artifact Reduction in CT. *IEEE Transactions on Medical Imaging*, 43(10):3521–3532, 2024. doi: 10.1109/TMI.2024.3416398.
- A. Kazerouni, E. K. Aghdam, M. Heidari, R. Azad, M. Fayyaz, I. Hachililoglu, and D. Merhof. Diffusion Models in Medical Imaging: A Comprehensive Survey. *Medical Image Analysis*, 88:102846, 2023. doi: 10.1016/j.media.2023.102846.
- J. F. Keating, J. Werier, P. Blachut, H. Broekhuysse, R. N. Meek, and P. J. O'Brien. Early Fixation of the Vertically Unstable Pelvis: The Role of Iliosacral Screw Fixation of the Posterior Lesion. *Journal of Orthopaedic Trauma*, 13(2):107–113, 1999. doi: 10.1097/00005131-199902000-00007.
- J. J. Lee, S. L. Rosenbaum, A. Martusiewicz, S. A. Holcombe, S. C. Wang, and J. A. Goulet. Transsacral Screw Safe Zone Size by Sacral Segmentation Variations. *Journal of Orthopaedic Research*, 33(2): 277–282, 2015. doi: 10.1002/jor.22739.
- Z. Li, Q. Gao, Y. Wu, C. Niu, J. Zhang, M. Wang, G. Wang, and H. Shan. Quad-Net: Quad-Domain Network for CT Metal Artifact Reduction. *IEEE Transactions on Medical Imaging*, 43(5):1866–1879, 2024. doi: 10.1109/TMI.2024.3351722.
- W.-A. Lin, H. Liao, C. Peng, X. Sun, J. Zhang, J. Luo, R. Chellappa, and S. K. Zhou. DuDoNet: Dual Domain Network for CT Metal Artifact Reduction. In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10504–10513, 2019. doi: 10.1109/CVPR.2019.01076.
- J. Liu, H. Li, B. Zeng, H. Wang, R. Kikinis, L. Joskowicz, and X. Chen. An End-to-End Geometry-Based Pipeline for Automatic Preoperative Surgical Planning of Pelvic Fracture Reduction and Fixation. *IEEE Transactions on Medical Imaging*, 44(1):79–91, 2025. doi: 10.1109/TMI.2024.3429403.
- P. Liu, H. Han, Y. Du, H. Zhu, Y. Li, F. Gu, H. Xiao, J. Li, C. Zhao, L. Xiao, X. Wu, and S. K. Zhou. Deep Learning to Segment Pelvic Bones: Large-Scale CT Datasets and Baseline Models. *International Journal of Computer Assisted Radiology and Surgery*, 16(5):749–756, 2021a. doi: 10.1007/s11548-021-02363-8.
- Y. Liu, E. Sangineto, W. Bi, N. Sebe, B. Lepri, and M. Nadai. Efficient Training of Visual Transformers with Small Datasets. In M. Ranzato, A. Beygelzimer, Y. Dauphin, P. S. Liang, and J. W. Vaughan, editors, *Advances in Neural Information Processing Systems*, volume 34, pages 23818–23830. Curran Associates, Inc., 2021b.
- A. Lugmayr, M. Danelljan, A. Romero, F. Yu, R. Timofte, and L. Van Gool. RePaint: Inpainting Using Denoising Diffusion Probabilistic Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11461–11471, 2022.
- Y. Lyu, W.-A. Lin, H. Liao, J. Lu, and S. K. Zhou. Encoding Metal Mask Projection for Metal Artifact Reduction in Computed Tomography. In *Medical Image Computing and Computer Assisted Intervention — MICCAI 2020*, pages 147–157. Springer International Publishing, 2020. ISBN 9783030597139. doi: 10.1007/978-3-030-59713-9_15.
- R. Macháček, L. Mozaffari, Z. Sepasdar, S. Parasa, P. Halvorsen, M. A. Riegler, and V. Thambawita. Mask-Conditioned Latent Diffusion for Generating Gastrointestinal Polyp Images. In *Proceedings of the 4th ACM Workshop on Intelligent Cross-Data Analysis and Retrieval, ICDAR '23*, pages 1–9, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9789400701863. doi: 10.1145/3592571.3592978.
- J. M. Matta and P. Tornetta, 3rd. Internal Fixation of Unstable Pelvic Ring Injuries. *Clinical Orthopaedics and Related Research*, (329):129–140, 1996. doi: 10.1097/00003086-199608000-00016.
- D. A. McLaren, G. A. Busel, H. R. Parikh, A. Only, J. Patterson, B. T. Gaston, R. McLemore, and B. Cunninghamham. Corridor-Diameter-Dependent Angular Tolerance for Safe Transiliac Sacral Screw Placement: An Anatomic Study of 433 Pelves. *European Journal of Orthopaedic Surgery & Traumatology*, 31(7): 1485–1492, 2021. doi: 10.1007/s00590-021-02913-5.
- T. Mendel, H. Noser, D. Wohlrab, K. Stock, and F. Radetzi. The Lateral Sacral Triangle—A Decision Support for Secure Transverse Sacroiliac Screw Insertion. *Injury*, 42(10):1164–1170, 2011. doi: 10.1016/j.injury.2010.03.016.
- A. N. Miller and M. L. C. Routt, Jr. Variations in Sacral Morphology and Implications for Iliosacral Screw Fixation. *The Journal of the American Academy of Orthopaedic Surgeons*, 20(1):8–16, 2012. doi: 10.5435/JAAOS-20-01-008.
- S. K. Mirza, G. C. Wiggins, C. Kuntz, J. E. York, C. Bellabarba, M. A. Knonodi, J. R. Chapman, and C. I. Shaffrey. Accuracy of Thoracic Vertebral Body Screw Placement Using Standard Fluoroscopy, Fluoroscopic Image Guidance, and Computed Tomographic Image Guidance: A Cadaver Study. *Spine*, 28(4):402–413, 2003. doi: 10.1097/01.BRS.0000048461.51308.CD.
- B. R. Moed and B. L. Geer. S2 Iliosacral Screw Fixation for Disruptions of the Posterior Pelvic Ring: A Report of 49 Cases. *Journal of Orthopaedic Trauma*, 20(6):378–383, 2006. doi: 10.1097/00005131-200607000-00002.
- H. R. Mostafavi and P. Tornetta, 3rd. Radiologic Evaluation of the Pelvis. *Clinical Orthopaedics and Related Research*, (329):6–14, 1996. doi: 10.1097/00003086-199608000-00003.
- A. Q. Nichol and P. Dhariwal. Improved Denoising Diffusion Probabilistic Models. In M. Meila and T. Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 8162–8171. PMLR, 2021.
- F. K. Noojin, A. L. Malkani, L. Haikal, C. Lundquist, and M. J. Voor. Cross-Sectional Geometry of the Sacral Ala for Safe Insertion of Iliosacral Lag Screws: A Computed Tomography Model. *Journal of Orthopaedic Trauma*, 14(1):31–35, 2000. doi: 10.1097/00005131-200001000-00007.
- H. S. Park, Y. E. Chung, and J. K. Seo. Computed Tomographic Beam-Hardening Artefacts: Mathematical Characterization and Analysis. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 373(2043):20140388, 2015. doi: 10.1098/rsta.2014.0388.
- W. Peebles and S. Xie. Scalable Diffusion Models with Transformers. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4172–4182, Los Alamitos, CA, USA, 2023. IEEE Computer Society. doi: 10.1109/ICCV51070.2023.00387.
- N. Peters, E. Haneda, J. Zhang, G. Karageorgos, W. Xia, J. Verburg, G. Wang, H. Paganetti, and B. De Man. A Hybrid Training Database and Evaluation Benchmark for Assessing Metal Artifact Reduction Methods for X-ray CT Imaging. *Medical Physics*, 52(10):e70020, 2025. doi: 10.1002/mp.70020.
- S. Radzi, G. Cowin, M. Robinson, J. Pratap, A. Volp, M. A. Schuetz, and B. Schmutz. Metal Artifacts from Titanium and Steel Screws in CT, 1.5T and 3T MR Images of the Tibial Pilon: A Quantitative Assessment in 3D. *Quantitative Imaging in Medicine and Surgery*, 4(3):163–172, 2014. doi: 10.3978/j.issn.2223-4292.2014.03.06.
- N. Ramadanov and S. Zabler. Ramadanov–Zabler Safe Zone for Sacroiliac Screw Placement: A CT-Based Computational Pilot Study. *Journal of Clinical Medicine*, 14(10):3567, 2025. doi: 10.3390/jcm14103567.
- F. Ramzan, Y. Kiberu, N. Jathanna, S. Jamil-Copley, R. H. Clayton, and C. Chen. CLAIM: Clinically-Guided LGE Augmentation for Realistic and Diverse Myocardial Scar Synthesis and Segmentation. In D. Cafolla, T. Rittman, and H. Ni, editors, *Artificial Intelligence in Healthcare*, pages 279–292, Cham, 2026. Springer Nature Switzerland. ISBN 978-3-032-00652-3. doi: 10.1007/978-3-032-00652-3_20.
- M. C. Reilly, C. M. Bono, B. Litkouhi, M. Sirkin, and J. F. Behrens. The Effect of Sacral Fracture Malreduction on the Safe Placement of Iliosacral Screws. *Journal of Orthopaedic Trauma*, 17(2): 88–94, 2003. doi: 10.1097/00005131-200302000-00002.
- L. Ren, A. Ferrero, and C. Favazza. Framework of Metal Insertion in the Projection Domain for Image

- Quality Optimization in Interventional Computed Tomography. *Journal of Medical Imaging*, 9(3): 035001, 2022. doi: 10.1117/1.JMI.9.3.035001.
- R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-Resolution Image Synthesis with Latent Diffusion Models. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10674–10685, 2022. doi: 10.1109/CVPR52688.2022.01042.
- O. Ronneberger, P. Fischer, and T. Brox. U-Net: Convolutional Networks for Biomedical Image Segmentation. In N. Navab, J. Hornegger, W. M. Wells, and A. F. Frangi, editors, *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, pages 234–241, Cham, 2015. Springer International Publishing. ISBN 978-3-319-24574-4.
- M. L. Rouitt, Jr., P. T. Simonian, and W. J. Mills. Iliosacral Screw Fixation: Early Complications of the Percutaneous Technique. *Journal of Orthopaedic Trauma*, 11(8):584–589, 1997a. doi: 10.1097/00005131-199711000-00007.
- M. L. C. Rouitt, P. T. Simonian, and J. Inaba. Iliosacral Screw Complications. *Operative Techniques in Orthopaedics*, 7(3):206–220, 1997b. doi: 10.1016/S1048-6666(97)80060-2. Pelvic Trauma: An update.
- S. C. Sayres, Y. Gu, S. Kiernan, B. A. DeSandis, A. J. Elliott, and M. J. O’Malley. Comparison of Rates of Union and Hardware Removal Between Large and Small Cannulated Screws for Calcaneal Osteotomy. *Foot & Ankle International*, 36(1):32–36, 2015. doi: 10.1177/1071100714549191.
- M. Selles, M. Maas, R. H. H. Wellenberg, D. J. Slotman, J. A. C. van Osch, I. M. Nijholt, and M. F. Boomsma. Is AI the Way Forward for Reducing Metal Artifacts in CT? Development of a Generic Deep Learning-Based Method and Initial Evaluation in Patients with Sacroiliac Joint Implants. *European Journal of Radiology*, 163, 2023. doi: 10.1016/j.ejrad.2023.110844.
- M. Selles, J. A. C. van Osch, M. Maas, M. F. Boomsma, and R. H. H. Wellenberg. Advances in Metal Artifact Reduction in CT Images: A Review of Traditional and Novel Metal Artifact Reduction Techniques. *European Journal of Radiology*, 170:111276, 2024. doi: 10.1016/j.ejrad.2023.111276.
- K. Singhrao, M. Zubair, T. Nano, J. E. Scholey, and M. Descovich. End-to-End Validation of Fiducial Tracking Accuracy in Robotic Radiosurgery Using MRI-Only Simulation Imaging. *Medical Physics*, 51(1):31–41, 2024. doi: 10.1002/mp.16857.
- H. E. Smith, P. S. Yuan, R. Sasso, S. Papadopolous, and A. R. Vaccaro. An Evaluation of Image-Guided Technologies in the Placement of Percutaneous Iliosacral Screws. *Spine*, 31(2):234–238, 2006. doi: 10.1097/01.brs.0000194788.45002.1b.
- J. Song, C. Meng, and S. Ermon. Denoising Diffusion Implicit Models. In *International Conference on Learning Representations*, 2021.
- Y. Song, T. Yao, S. Peng, M. Zhu, M. Meng, J. Ma, D. Zeng, J. Huang, Z. Bian, and Y. Wang. b-MAR: Bidirectional Artifact Representations Learning Framework for Metal Artifact Reduction in Dental CBCT. *Physics in Medicine & Biology*, 69(14):145010, 2024. doi: 10.1088/1361-6560/ad3c0a.
- Cannulated Screws 3.0/3.5/4.0/4.5/6.5/7.0/7.3. Surgical Technique*. Synthes GmbH, Oberdorf, Switzerland, 2006. Document 036.000.094.
- Synthes Screw Reference Chart: Screws, Drill Bits, Taps and Guide Wires*. Synthes (USA), Paoli, PA, USA, 2002. Print code GP0644-C 4/04 J1364-C.
- 6.5 mm and 7.3 mm Cannulated Screws: Technique Guide*. Synthes (USA), West Chester, PA, USA, 2003. Document J4552-E, printing 2/10.
- N. C. Tejwani, D. Raskolnikov, T. McLaurin, and R. Takemoto. The Role of Computed Tomography for Postoperative Evaluation of Percutaneous Sacroiliac Screw Fixation and Description of a “Safe Zone”. *American Journal of Orthopedics*, 43(11):513–516, 2014.
- D. Templeman, A. Schmidt, J. Freese, and I. Weisman. Proximity of Iliosacral Screws to Neurovascular Structures After Internal Fixation. *Clinical Orthopaedics and Related Research*, (329):194–198, 1996. doi: 10.1097/00003086-199608000-00023.
- A. R. Vaccaro, S. J. Rizzolo, T. J. Allardyce, M. Ramsey, J. Salvo, R. A. Balderston, and J. M. Cotler. Placement of Pedicle Screws in the Thoracic Spine. Part I: Morphometric Analysis of the Thoracic Vertebrae. *The Journal of Bone and Joint Surgery. American Volume*, 77(8):1193–1199, 1995a. doi: 10.2106/00004623-199508000-00008.
- A. R. Vaccaro, S. J. Rizzolo, R. A. Balderston, T. J. Allardyce, S. R. Garfin, C. Dolinskas, and H. S. An. Placement of Pedicle Screws in the Thoracic Spine. Part II: An Anatomical and Radiographic Assessment. *The Journal of Bone and Joint Surgery. American Volume*, 77(8):1200–1206, 1995b. doi: 10.2106/00004623-199508000-00009.
- H. J. P. van Bosse, D. Lee, E. R. Henderson, D. A. Sala, and D. S. Feldman. Pelvic Positioning Creates Error in CT Acetabular Measurements. *Clinical Orthopaedics and Related Research*, 469(6): 1683–1691, 2011. doi: 10.1007/s11999-011-1827-9.
- E. W. van den Bosch, C. M. A. van Zwienen, and A. B. van Vugt. Fluoroscopic Positioning of Sacroiliac Screws in 88 Patients. *The Journal of Trauma*, 53(1):44–48, 2002. doi: 10.1097/00005373-200207000-00009.
- D. Wagner, L. Kamer, T. Sawaguchi, R. Geoff Richards, H. Noser, M. Uesugi, C. Ossendorf, and P. M. Rommens. Critical Dimensions of Trans-Sacral Corridors Assessed by 3D CT Models: Relevance for Implant Positioning in Fractures of the Sacrum. *Journal of Orthopaedic Research*, 35(11):2577–2584, 2017. doi: 10.1002/jor.23554.
- H. Wang, Y. Li, D. Meng, and Y. Zheng. Adaptive Convolutional Dictionary Network for CT Metal Artifact Reduction. In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence, IJCAI-2022*, pages 1401–1407. International Joint Conferences on Artificial Intelligence Organization, 2022. doi: 10.24963/ijcai.2022/195.
- H. Wang, Y. Wu, Y. Wang, D. Wei, X. Wu, J. Ma, and Y. Zheng. Adaptive Weighting Based Metal Artifact Reduction in CT Images. *IEEE Transactions on Medical Imaging*, 44(6):2408–2423, 2025. doi: 10.1109/TMI.2025.3534316.
- Z. Wang, C. Vandersteen, T. Demarcy, D. Gnansia, C. Raffaelli, N. Guevara, and H. Delingette. Deep Learning Based Metal Artifacts Reduction in Post-operative Cochlear Implant CT Imaging. In D. Shen, T. Liu, T. M. Peters, L. H. Staib, C. Essert, S. Zhou, P.-T. Yap, and A. Khan, editors, *Medical Image Computing and Computer Assisted Intervention – MICCAI 2019*, pages 121–129, Cham, 2019. Springer International Publishing. ISBN 978-3-030-32226-7. doi: 10.1007/978-3-030-32226-7_14.
- J. Wasserthal, H.-C. Breit, M. T. Meyer, M. Pradella, D. Hinck, A. W. Sauter, T. Heye, D. T. Boll, J. Cyriac, S. Yang, M. Bach, and M. Segeroth. TotalSegmentor: Robust Segmentation of 104 Anatomic Structures in CT Images. *Radiology: Artificial Intelligence*, 5(5):e230024, 2023. doi: 10.1148/ryai.230024.
- L.-P. Wu, Y.-K. Li, Y.-M. Li, Y.-Q. Zhang, and S.-Z. Zhong. Variable Morphology of the Sacrum in a Chinese Population. *Clinical Anatomy*, 22(5):619–626, 2009. doi: 10.1002/ca.20809.
- M. Wu, P. FitzGerald, J. Zhang, W. P. Segars, H. Yu, Y. Xu, and B. De Man. XCIST — An Open Access X-ray/CT Simulation Toolkit. *Physics in Medicine & Biology*, 67(19):194002, 2022. doi: 10.1088/1361-6560/ac9174.
- K. Xie, L. Gao, Y. Zhang, H. Zhang, J. Sun, T. Lin, J. Sui, and X. Ni. Metal Implant Segmentation in CT Images Based on Diffusion Model. *BMC Medical Imaging*, 24(1):204, 2024. doi: 10.1186/s12880-024-01379-1.
- R. Xu, N. A. Ebraheim, K. Douglas, and R. A. Yeasting. The Projection of the Lateral Sacral Mass on the Outer Table of the Posterior Ilium. *Spine*, 21(7):790–794, 1996. doi: 10.1097/00007632-199604010-00002.
- M. Yazdi, M. A. Lari, G. Bernier, and L. Beaulieu. An Opposite View Data Replacement Approach for Reducing Artifacts Due to Metallic Dental Objects. *Medical Physics*, 38(4):2275–2281, 2011. doi: 10.1118/1.3566016.
- P. L. Yeap, X. Du, M. Zhou, A. Hoole, G. C. Barnett, and R. Jena. Few-Shot CBCT-Based Synthetic CT Generation with Denoising Diffusion Probabilistic Model. *Medical Physics*, 52(11):e70126, 2025. doi: 10.1002/mp.70126.
- L. Yu, Z. Zhang, X. Li, and L. Xing. Deep Sinogram Completion With Image Prior for Metal Artifact Reduction in CT Images. *IEEE Transactions on Medical Imaging*, 40(1):228–238, 2021. doi: 10.1109/TMI.2020.3025064.
- S. Yun, S. Hyun, D.-I. Choi, and S. Cho. A Strategy for Simulation-Driven CT Metal Artifact Reduction Toward Improving Network Generalizability. *Medical Physics*, 53(2):e70336, 2026. doi: 10.1002/mp.70336.
- H. Zhang, Y. Liu, J. Yang, S. Wan, X. Wang, W. Peng, and P. Fua. LeFusion: Controllable Pathology Synthesis via Lesion-Focused Diffusion Models. In Y. Yue, A. Garg, N. Peng, F. Sha, and R. Yu, editors, *International Conference on Learning Representations*, volume 2025, pages 13232–13253, 2025a.
- L. Zhang, A. Rao, and M. Agrawala. Adding Conditional Control to Text-to-Image Diffusion Models. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3813–3824, 2023. doi: 10.1109/ICCV51070.2023.00355.
- Y. Zhang and H. Yu. Convolutional Neural Network Based Metal Artifact Reduction in X-Ray Computed Tomography. *IEEE Transactions on Medical Imaging*, 37(6):1370–1381, 2018. doi: 10.1109/TMI.2018.2823083.
- Y. Zhang, H. Luo, D. He, Z. Guan, J. Zhao, and Z. Yang. Enhancing Intraoperative Pedicle Screw Planning Accuracy via Diffusion-Based Synthetic CT. *Journal of X-Ray Science and Technology*, 34(4):525–540, 2026. doi: 10.1177/08953996261443500.
- Z. Zhang, L. Yao, B. Wang, D. Jha, G. Durak, E. Keles, A. Medetalibeyoglu, and U. Bagci. DiffBoost: Enhancing Medical Image Segmentation via Text-Guided Diffusion Model. *IEEE Transactions on Medical Imaging*, 44(9):3670–3682, 2025b. doi: 10.1109/TMI.2024.3519307.
- Y. Zhao, J. Li, D. Wang, and W. Lian. Parameters of Lengthened Sacroiliac Screw Fixation: A Radiological Anatomy Study. *European Spine Journal*, 21(9):1807–1814, 2012. doi: 10.1007/s00586-012-2367-z.
- Y. Zhu, H. Zhao, T. Wang, L. Deng, Y. Yang, Y. Jiang, N. Li, Y. Chan, J. Dai, C. Zhang, Y. Li, Y. Xie, and X. Liang. Sinogram Domain Metal Artifact Correction of CT via Deep Learning. *Computers in Biology and Medicine*, 155:106710, 2023. doi: 10.1016/j.combiomed.2023.106710.
- Z. Zhu and F. Liao. Optimal Position and Apex-Subchondral Bone Distance of Cannulated Screws for Internal Fixation of Femoral Neck Fracture: A Radiological Study. *Journal of Orthopaedic Research*, 41(7):1546–1554, 2023. doi: 10.1002/jor.25503.
- B. H. Ziran, W. R. Smith, J. Towers, and S. J. Morgan. Iliosacral Screw Fixation of the Posterior Pelvic Ring Using Local Anaesthesia and Computerised Tomography. *The Journal of Bone and Joint Surgery. British Volume*, 85(3):411–418, 2003. doi: 10.1302/0301-620x.85b3.13119.
- B. H. Ziran, A. D. Wasan, D. M. Marks, S. A. Olson, and M. W. Chapman. Fluoroscopic Imaging Guides of the Posterior Pelvis Pertaining to Iliosacral Screw Placement. *The Journal of Trauma*, 62(2):347–356, 2007. doi: 10.1097/01.ta.0000229789.18211.85.
- J. Zwingmann, G. Konrad, E. Kotter, N. P. Südkamp, and M. Oberst. Computer-Navigated Iliosacral Screw Insertion Reduces Malposition Rate and Radiation Exposure. *Clinical Orthopaedics and Related Research*, 467(7):1833–1838, 2009. doi: 10.1007/s11999-008-0632-6.
- J. Zwingmann, G. Konrad, A. T. Mehlhorn, N. P. Südkamp, and M. Oberst. Percutaneous Iliosacral Screw Insertion: Malpositioning and Revision Rate of Screws with Regards to Application Technique (Navigated vs. Conventional). *The Journal of Trauma*, 69(6):1501–1506, 2010. doi: 10.1097/TA.0b013e3181d862db.
- J. Zwingmann, O. Hauschild, G. Bode, N. P. Südkamp, and H. Schmal. Malposition and Revision Rates of Different Imaging Modalities for Percutaneous Iliosacral Screw Fixation Following Pelvic Fractures: A Systematic Review and Meta-Analysis. *Archives of Orthopaedic and Trauma Surgery*, 133(9): 1257–1265, 2013. doi: 10.1007/s00402-013-1788-4.