

MetalSynth-Pelvis: Conditioned synthesis of osteosynthesis implants and local metal artifacts in pelvic CT scan volumes using Multi-Window Latent Diffusion

Kiara Alexandra Balcázar Santa Cruz — UTEC · Advisor: Víctor Flores Benites

1. Problem Statement

Automatic pelvic computed tomography (CT) analysis and volumetric segmentation depend on sufficiently annotated data. Metal artifacts can obscure anatomical boundaries near implants, motivating evaluation of segmentation robustness in this region. CTPelvic1K (Liu et al., 2021) includes a CLINIC-metal subset whose limited bone annotation constrains supervised evaluation. Its name does not establish that every case contains pelvic osteosynthesis, nor that other subsets are free of metal or foreign objects. The local audit will therefore determine implant types, eligible training anatomy and annotated evaluation cases through three-dimensional and multiplanar review. The magnitude of peri-implant segmentation degradation must be measured in this study rather than assumed from the dataset description.

While generative models based on diffusion processes excel at synthesizing visually plausible medical anomalies, they structurally fail when applied to metallic implants. State-of-the-art diffusion-based lesion synthesis models fall into two families, and neither represents the intensity changes that the synthesized object induces outside itself. Mask-bounded inpainting methods condition generation on a lesion mask: CLAIM (Ramzan et al., 2026) recomposes the background from the original image at every denoising step, DiffTumor (Chen et al., 2024) models nothing outside the mask, and LGESynthNet (Jacob et al., 2026) is formulated as inpainting, although its latent compression still perturbs the whole slice incidentally rather than by design. Globally conditioned methods such as DiffBoost (Zhang et al., 2025) instead synthesize the entire slice from noise, using the mask only as an edge constraint. None of these works states an assumption of tissue non-rigidity; what they share is that no mechanism, and no training signal, accounts for object-induced intensity changes beyond the mask. This is physically incompatible with metal implants, which are rigidly geometric and produce severe attenuation profiles (beam hardening, photon starvation, and high-amplitude streaking) that propagate beyond the implant's boundaries and degrade the signal of the surrounding tissue; reduction networks describe them as present globally in the reconstructed image (Li et al., 2024). De Man et al. (1999) attribute such streaks to beam hardening, scatter, noise and the exponential edge-gradient effect, and show them radiating from the metal along the directions of highest attenuation; their simulation, however, uses an iron rod and dental amalgam in a two-dimensional fan-beam geometry, excludes non-linear partial volume, and reports no spatial extent, so the size of the affected peri-implant band remains unmeasured in the literature.

Physics-based metal artifact simulation provides a reference for appearance generation. This thesis adopts the hybrid simulation protocol of Peters et al. (2025), implemented with CatSim within XCIST (De Man et al., 2007; Wu et al., 2022), as its physical comparison protocol. Their validation compares simulations against a physical phantom scanned on a commercial scanner; it does not cover the hybrid step (clinical images as input, with an empirical frequency-boosting compensation), the vendor-neutral geometry used to generate their data, or pelvic osteosynthesis in three dimensions, and the paper does not state whether patient and metal are projected jointly, which this work verifies in the simulator code before reproduction. Their training data use randomized metal placement, whereas their scoring benchmark uses expert placement. The unresolved placement problem is therefore scalable, anatomically constrained placement, rather than scalable metal insertion itself.

Even if automated placement were achieved, mere anatomical containment is insufficient; a valid simulation must reflect true surgical admissibility. Traditional deterministic surgical planning pipelines (Liu et al., 2025) attempt to optimize a single "perfect" trajectory. Yet real-world clinical data reveal a different reality. Grading iliosacral screws on a four-level cortical-breach scale, Zwingmann et al. (2009) report two distinct postoperative-CT distributions: grade 0 in 69% of navigated placements (grades 1–3: 15%, 8%, 8%) against grade 0 in 40% of conventional placements (37%, 11.5%, 11.5%). That scale reaches iliosacral fixation through Smith et al. (2006), who apply it there, but it originates in the pedicle-screw literature: Gertzbein and Robbins (1990) tabulate canal encroachment from the medial pedicle border in six 2 mm bins rather than four numbered grades, anchor the 2 mm cut to the epidural space, extrapolate the 4 mm cut from a single myelographic CT, and state no justification for the 6 and 8 mm cuts. The four-grade form (0, <2, 2–4, >4 mm) appears literally in Mirza et al. (2003), who declare the thresholds taken from prior studies, measured perforation in cadaveric thoracic spines rather than on CT, and caution that these "safe zone" thresholds do not apply to the thoracic spine and are likely different for different directions of screw perforation. The 2 mm bin width is therefore adopted here as a geometric convention, as with the 10 mm corridor criterion below, and not as a validated clinical threshold; its transfer from the spinal canal to the iliosacral corridor is an explicit assumption of this work. These are two surgical-technique populations that differ significantly, not one aggregate range; this work therefore targets the ordinal distributions themselves rather than a single imperfection rate. The distinction is not merely presentational. In a meta-analysis by the same first author, Zwingmann et al. (2013) pool malposition rates of 2.6% for conventional placement against 0.1% under CT navigation, and the 2009 cohort enters that pool with zero events in both arms, because the included series record malposition only where a revision followed; the authors themselves note that no comparable grading systems are given across studies. Aggregate malposition rates and ordinal cortical-breach grades therefore measure different constructs, and only the latter is a usable target for a placement sampler. Sacral level also warrants separate treatment on anatomical grounds. Using multiplanar CT measurements, Gardner et al. (2010) report mean minimum iliosacral cross-sectional areas at S1 and S2 of 346 and 109.3 mm² in normal sacra, compared with 222 and 220.1 mm² in dysmorphic sacra. The approximately equal areas in the latter group support measuring each level on the individual anatomy rather than assuming a universal corridor-size hierarchy. These measurements concern an uninjured side and ideal central trajectories; they do not describe clinical malposition frequencies. In a retrospective clinical series, van den Bosch et al. (2002) reported postoperative neurologic complaints in 6 of 31 patients with the lower screw in S2, versus 1 of 49 with both screws in S1 ($p = 0.01$). These are patient-level outcomes by screw configuration, not screw-malposition rates by level; the authors also identify a learning-curve bias associated with the change in placement practice. This clinical association supports retaining S1/S2 as a sampler variable, but supplies neither a causal level effect nor an ordinal cortical-breach prior for S2. Cortical breach by level need not follow the same ordering as patient-level neurologic outcome: on postoperative CT, Herman et al. (2017) report breach in 36.5% of S1 screws against 14.8% at S2 ($p = 0.035$), with 32% of screws breaching overall and no screw lying entirely outside bone. The two findings are measured on different quantities and are not in conflict; the latter bounds the range of malpositions the sampler must be able to produce. The ordinal benchmark from Zwingmann et al. (2009) applies to S1: the sacral level is explicit for the conventional arm and not reported for the navigated arm, where it is assumed to be S1, and grades are counted per screw, some patients having received more than one; S2 placements are evaluated descriptively for geometry and breach grades without claiming clinical distribution matching. Its grades were assigned on postoperative CT by one independent radiologist, with no slice thickness, reconstruction interval or inter-observer agreement reported, and the text does not state to which grade a perforation of exactly 2 or 4 mm belongs. Detection itself may depend on image resolution: Tejwani et al. (2014) describe more foraminal penetration on 2.5 mm

than on 5.0 mm sections, although at $P = .3$. Matching against this benchmark is therefore read as agreement between grade distributions measured at an unknown reference resolution, not as voxel-level equivalence, and the grade boundaries used here are a declared convention. The premise that no canonical pose exists is not only inferred from surgical error rates; it is supported anatomically. Measuring the orientation of the bony surfaces that bound the corridor across seventeen cadaveric pelvises, Ziran et al. (2007) report inter-specimen coefficients of variation of 7–25% for most surfaces, rising to 43%, 47–52% and, for the orientation of the superior S1 ala in the sacral frontal plane, 97–140%. A corridor that varies this much between individuals cannot be served by one fixed trajectory. Training a diagnostic network solely on "perfect" deterministic trajectories will inevitably cause it to fail when encountering the true variance of clinical errors. Therefore, a generative sampler must not optimize for perfection, but instead sample from the measured ordinal distribution of clinical malpositions, bounded by pelvic density representations (Arand et al., 2019) and by osseous-corridor viability. For the latter, McLaren et al. (2021) supply the operational element that a qualitative safe-zone description (Ramadanov and Zabler, 2025) does not: a reproducible corridor-diameter procedure over 433 pelvises, yielding transiliosacral angular tolerances of $1.53 \pm 0.57^\circ$ at S1 and $1.02 \pm 0.33^\circ$ at S2. Their ≥ 10 mm corridor criterion is not an empirically validated safety threshold and is not treated as one here: Kaiser et al. (2014), from whom it is taken, describe it as a chosen conservative value and justify it dimensionally, as clearance of 1–2 mm around a 6.3–8 mm screw. It is therefore used as a geometric clearance convention rather than as an inherited clinical standard. Kaiser et al. (2014) additionally supply what a corridor diameter alone does not: a reference frame computable on a CT volume, obtained by reformatting along the sacral axis perpendicular to the superior end plate of S1, with corridor angulation measured coronally against the line joining the iliac crests and axially against the line joining the posterior iliac spines, and a usable-length rule requiring at least 5 mm of clearance to the cortex on either side. The pose distribution sampled within that frame remains an explicit construction of this work.

Recognized Gap. Current attempts to bridge the metal-artifact domain shift are bottlenecked by three mutually exclusive paradigms: physical simulators do not supply scalable anatomically constrained placement, biological diffusion models cannot generate rigid global streaks, and deterministic surgical planners ignore the statistical reality of clinical malposition. To the best of the author's knowledge, there is no diffusion-based method that overcomes these barriers by combining: **(C1)** Non-biological rigid implant geometries, which avoid the threshold-dependent size distortion of fixed-HU metal segmentation (over-coverage in simulated slices, Xie et al., 2024; fragmentation of pelvic screws in the local audit); **(C2)** A surgically constrained placement sampler that accurately emulates the population distribution of clinical malpositions; and **(C3)** The transfer of multi-window HU encoding from artifact *removal*, where it originates, to artifact *synthesis*, coupled with a generation band (B_δ) extended beyond the implant mask so that streaking and beam hardening are modelled in a bounded region outside the metal itself. The multi-window framework is not claimed as novel here: it is established in metal-artifact-reduction work, where published uses apply the windows to per-window reconstruction stages or to the training loss rather than to the encoding of the input (Wang et al., 2025; Li et al., 2024), and Wang et al. (2025) attribute it in turn to earlier multiple-window learning. What is claimed is its use for generation, together with B_δ , for which no published precedent quantifying a peri-implant band was found. This is the gap this work aims to fill.

Filling it rests on one assumption that the present literature neither grants nor refutes: that the *appearance* of a metal artifact can be generated in the image domain, without access to projection data. The mechanisms involved are non-linear and are manipulated in the sinogram (De Man et al., 1999), and dual-domain reduction networks report a penalty for operating on images alone: Lin et al. (2019) obtain 31.45 dB PSNR from their image-domain variant against 33.51 dB for the full dual-domain network. That image-domain variant is fed a reconstruction of an already-interpolated sinogram, but the penalty does not rest on this detail: Li et al. (2024), whose image-only branch starts from the metal-corrupted reconstruction, also report it below their multi-domain network and state that working in a single domain cannot effectively recover the true tissues. One qualification keeps these results from settling the question: both address artifact *removal*, which Lin et al. (2019) describe as ill-posed because information has been destroyed, whereas synthesis carries no such recovery requirement. This work therefore treats image-domain generability not as a premise but as part of what its distributional physical-coherence evaluation puts to the test, with the physics-based protocol of Peters et al. (2025) supplying the mechanism-faithful reference against which the generated appearance is judged.

2. Research Question

What improvement does a surgical admissibility-constrained (bone density and cortical containment) 3D placement sampler, a multi-window Latent Diffusion Model (LDM) conditioned on the rigid implant geometries and an extended generation band (B_δ), have in terms of distributional physical-coherence of the synthesized pelvic osteosynthesis implants and their local artifacts, when compared with both naive copy-paste insertion and the adopted physics-based simulation protocol of Peters et al. (2025)? Downstream segmentation impact is deliberately excluded from this study and is stated as out of scope below.

3. Hypothesis

The synthetic generation of medical anomalies and metallic implants using a rigid geometry and extended spatial boundaries, HU-conditioned diffusion model and whose placement is surgically enforced via a density-based placement sampler is more physically-coherent than unconstrained generative baseline methods and on the same footing as the adopted explicit physics-based simulation protocol (Peters et al., 2025). Placement coherence is quantified by Surgical Admissibility of Placement (SAP), the single metric introduced by this work, which emulates the measured ordinal distribution of real-world implant placement. Appearance coherence is quantified with the metrics of the adopted protocol under their published names, not with newly coined ones.

4. Objectives

General Objective. Design, implement and evaluate a physics-and-anatomy-aware latent diffusion pipeline with rigid surgical constraints and extended HU-based conditioning, capable of generating physically coherent metallic implants intended, in future work, to improve the robustness of downstream segmentation networks.

1. **Validation of a multi-window representation (Compulsory Go/No-Go):** Ensure HU reconstruction error (HU to multi-window to VAE to HU again) within the bone tissue, delimited by the same 150 HU bone threshold used by the adopted evaluation protocol (Peters et al., 2025; Haneda et al., 2025), below a 25 HU threshold (Mean Absolute Error (MAE)). No published pass/fail threshold on CT-number accuracy was found, so this criterion is not taken from a source: it is anchored on the error level that current metal-artifact methods reach on the same scale. The NMAR algorithm that calibrates the community benchmark reports an RMSE of 20.2 HU, an image-domain diffusion model 12.3 HU (Karageorgos et al., 2024), and a latent diffusion model 12.74 HU (Yun et al., 2026). Because those values are RMSE while

this criterion is MAE, and RMSE is never below MAE for the same residuals, the comparison fixes the order of magnitude rather than an equivalence; this is declared rather than resolved. A concurrent preprint reports that frozen video autoencoders transfer to CT without medical fine-tuning (Chen et al., 2026); it reports PSNR, SSIM and MSE without declared units and no HU error, evaluates no bone or metal region, uses no pelvic or metal-bearing volumes, and declares intensity clipping to $[-1000, 1000]$ HU for its generation experiments, so it does not address the dense-bone and metal range this criterion targets. In volumes with metal, the same round-trip error is additionally reported inside the band B_δ around the metal, descriptively and outside the Go/No-Go: that band holds the low-amplitude streaks the synthesizer must render, and an error restricted to bone does not show whether the autoencoder preserves them, since the first stage of latent diffusion is designed to remove high-frequency detail (Rombach et al., 2022). The gate is evaluated on held-out patients for the pretrained autoencoder and for a decoder-adapted variant with the encoder frozen; if neither passes, Objective 3 is not pursued and the No-Go result is reported as the outcome. This will serve as a guarantee that the intensity gradients that drive downstream applications are preserved.

2. **Design of surgically constrained placement sampler:** propose position-orientation 3D pose of a parametric rigid screw model constrained by bone density read from the CT numbers of each volume, motivated by the systematic heterogeneity of pelvic bone mass (e.g., low grey values in the sacral ala) shown by a population-mean statistical model (Arand et al., 2019), cortical boundaries with a 5 mm clearance rule, and osseous-corridor viability (McLaren et al., 2021), expressed in the CT-computable reference frame of Kaiser et al. (2014) (sacral-axis reformat; coronal and axial corridor angulation against iliac-crest and posterior-iliac-spine landmarks). Corridor geometry is measured on anatomical masks from TotalSegmentator 2.18.0 (Wasserthal et al., 2023), which is built on nnU-Net (Isensee et al., 2021) (task `total`, default crop; robust crop as sensitivity analysis), quality-controlled per case and cleaned of connected components below 0.1% of their structure, rather than on an HU threshold. The cited publication evaluates a 104-structure model and reports no accuracy for the S1 label used here, so its reported accuracy is not assumed for the version and classes used in this work. The sampler must yield a distribution of placements emulating the two measured ordinal malposition distributions of Zwingmann et al. (2009), conditioned on surgical technique within S1, rather than a single artificial "optimized" trajectory. Sacral level remains a geometric conditioning variable, anatomically motivated by the segment-specific corridor measurements of Gardner et al. (2010); S2 results are reported descriptively, because the clinical comparison of van den Bosch et al. (2002) concerns neurologic complaints by screw configuration and does not provide an ordinal S2 placement distribution. Sacral-morphology phenotypes are deliberately not used as a stratifying variable: their reported effect on safe-zone size is not consistent across studies, and the sampler instead measures corridor geometry directly on each volume.
3. **Design of a conditioned synthesizer:** 2.5D ControlNet-guided Latent Diffusion Model on a Stable Diffusion 1.5 backbone. ControlNet presupposes a frozen pretrained base and its latent space (Zhang et al., 2023); the encoder is therefore kept frozen and only the decoder is adapted to the multi-window CT encoding, conditional on passing the Objective 1 gate. A crucial detail in the formulation is that the synthesis mask must be extended with a spatial band B_δ ($\sim 12\text{mm}$) over the implant boundary, to allow a more explicit manifestation of the streaking artifacts and beam hardening that occur outside of the implant mask in the generated image. The band width is a design parameter: it bounds where artifact is generated and deliberately truncates far-field streaks, whose spatial extent remains unmeasured in the reviewed literature.
4. **Formalization of the placement metric and adoption of an inherited appearance protocol:** Formulation of Surgical Admissibility of Placement (SAP) as the only metric introduced by this work, evaluated with Wasserstein-1 distances against the measured ordinal distributions. Appearance is evaluated with the metrics of the adopted protocol under their published names — bone integrity, metal integrity and streak amplitude (Peters et al., 2025) — and the previously proposed BFC and ISC are withdrawn, so that inherited measures are not renamed as contributions.

Explicitly out of scope. Downstream segmentation evaluation (Dice, HD95) is excluded from this study, and constraint-by-constraint ablations are deferred to future work. Two auditable conditions drive the exclusion rather than time alone: within the locally available cohort only a minority of the metal-bearing volumes carry verified bone annotation, and only part of the published collection is held locally, so a downstream result would rest on a cohort too small and too partially annotated to support the claim. The synthesis arms and the placement sampler are therefore evaluated on physical coherence and surgical admissibility alone; downstream utility is left as future work.

5. Expected Results

Goal	Metric(s)	Expected Outcome
Surgical Admissibility (Only novel metric)	SAP: four-level Cortical Breach Grade (Smith et al., 2006), with boundaries fixed by convention as grade 0 for no perforation, 1 for $(0, 2)$, 2 for $[2, 4]$ and 3 for $(4, \infty)$ mm; density-zone fraction; corridor viability (McLaren et al., 2021) in the reference frame of Kaiser et al. (2014).	Distribution matching: Wasserstein-1 distance shows generated placements approaching the two measured ordinal distributions of Zwingmann et al. (2009), conditioned on technique within S1. S2 geometry and breach grades are reported descriptively, without a clinical ordinal target.
Appearance Coherence (Inherited protocol, published names)	Bone integrity; metal integrity; streak amplitude (Peters et al., 2025). Inversion of each metric for synthesis is documented, and the 0–4 scale is not carried over.	Lower profile discrepancy than naive copy-paste insertion, with streak amplitudes statistically comparable to, or better than, the adopted physics-based protocol.
Representation Viability (Standard Encoding Eval)	MAE in HU (Round-trip; bone, and B_δ descriptively)	Error in the bone tissues strictly below 25 HU threshold, a level anchored on the CT-number accuracy reported for NMAR (20.2 HU) and for learned metal-artifact methods (12.3–12.74 HU) (Karageorgos et al., 2024; Yun et al., 2026), not on a published pass/fail criterion. Log-compressed metal ranges preserve the structural topology.
Preservation outside B_δ (No-damage check)	RMSE, SSIM against the original CT, restricted to the complement of B_δ .	No measurable change where none is intended; deviation inside B_δ is the intended signal and is not scored as error.

6. Datasets

Primary Dataset. CTPelvic1K (Liu et al., 2021) will be used as the source of patient anatomy and real artifact observations. The locally available dataset6 and dataset7 volumes are inventoried separately from the full published collection; the CLINIC-metal label alone does not establish implant type or annotation availability. Acquisition and reconstruction settings, including any vendor metal artifact reduction or monoenergetic reconstruction, are not reported in the dataset publication; since both change artifact appearance (Selles et al., 2024), real artifact observations are treated as an appearance reference of unknown reconstruction, not as a physical ground truth. *Strict Isolation Rule.* No patient may contribute to both training and evaluation. The implant geometries consumed by the proposed sampler and by the physical simulation arm are parametric and are not derived from any evaluation volume. Duplicate-volume screening and patient grouping are required to enforce this separation.

Implant geometry source. Implants are modelled as parametric rigid screws rather than extracted from the metal-bearing volumes. Shaft diameter is drawn from the calibres reported for iliosacral fixation, 6.5–8.0 mm (Gardner et al., 2010) and 6.3–8 mm (Kaiser et al., 2014). Screw length is not taken from a fixed published range: it is bounded on each volume by the measured corridor under the usable-length rule of Kaiser et al. (2014). Features without a textual source in the reviewed literature (thread, cannulation, head and washer) are simplified to a smooth cylinder, and this simplification is declared as a limitation. The real implants of the metal-bearing subset do not enter the geometry bank; they serve as a descriptive contrast for placements, as an appearance reference for synthesis, and as a test–retest control where the same implant was imaged twice. The local audit motivates this choice: screws extracted from these volumes with a fixed HU threshold appear fragmented, with shaft diameters below the published calibres, and the extracted geometry changes between acquisitions of the same implant.

Physical comparison protocol. The hybrid simulation and evaluation protocol of Peters et al. (2025), also associated with the AAPM CT-MAR Challenge (Haneda et al., 2025), is adopted instead of an independently validated reimplement of XCIST/CatSim. Challenge data remain excluded from model training, and the scoring benchmark is reserved for final evaluation. Reproduction will document the code revision, configuration, reconstruction settings and access conditions. The physical arm is run on a reduced evaluation subset of patients. Both synthesis arms will use the same source anatomy and implant poses; a no-metal simulation/reconstruction control will quantify changes introduced by the physical arm alone. The published protocol is two-dimensional and evaluates MAR: extending it to pelvic osteosynthesis synthesis requires an explicit adaptation and does not inherit its validation automatically. RMSE and SSIM against the original CT will assess preservation outside the generation band, where changes are not intended; artifact realism within the affected region requires separate evaluation. No unverified numerical acceptance thresholds or MAR score calibration are adopted here.

Field limitation: corridor geometry has only been characterized in intact pelves. This is not a limitation of the present cohort but of the available literature. Every source that publishes sacral corridor geometry measured it in pelves without implants and without fracture: Kaiser et al. (2014) excluded scans with implants obscuring the lumbosacral junction, McLaren et al. (2021) used uninjured scans, and Ziran et al. (2007) worked on cadaveric specimens and warn that fracture, displacement and soft-tissue injury would alter what is visible. The reference frame adopted here rests on landmarks at the lumbosacral junction, which is precisely where streaking concentrates, so whether those landmarks remain locatable in metal-bearing volumes is not answered by any published source. It was therefore quantified on the local cohort. Landmarks were located by a heuristic on each volume, and a landmark region was counted as contaminated when its dark-streak or metal fraction exceeded the maximum observed in 69 metal-free patients of the same cohort. Among 65 patients with orthopaedic metal, all five landmarks were found within the field of view in 57. After a clinician reviewer (one surgeon, otorhinolaryngology), blinded to the automatic audit, confirmed the S1 level on every patient, the frame was computable in 48 patients and free of contamination at all five landmark regions in 29. Losses were driven more by field-of-view truncation of the iliac crests (7 patients) and by landmark localisation than by artifact contamination; in the 17 patients with a screening-detected iliosacral screw, the S1 level was correct in all and its region was contaminated in 3. The frame is therefore computed on metal-bearing volumes where it is computable; the handling of field-of-view-truncated cases is assessed separately, and the heuristic localisation of the iliac crests and posterior iliac spines has not been clinically reviewed.

Corridor measurement on the local cohort. Among 103 patients without osteosynthesis, 91 had a located S1 and 72 passed level quality control. Level discordance excluded 11 of 34 patients with non-orthopaedic objects and 7 of 57 clean patients, and one further patient was excluded because S1 reached the field-of-view border; this imbalance is declared. The maximum transsacral corridor diameter had a median of 9.5 mm (IQR 7.4–11.7), and the 10 mm convention was met in 29 of 72 patients with the default crop and in 27 with the robust crop; the crop is therefore reported as a measurement uncertainty that reaches the binary viability outcome. Component cleaning changed neither the transsacral corridor diameter nor its 10 mm viability in this cohort at any tested fraction up to 5%; the fraction, like the occupancy rule below, was fixed after inspecting these results. Along the axis of the widest corridor, the median fraction of voxels at or below 150 HU inside the sacral masks was 0.42, so an HU-based bone definition would exclude much of the corridor. In metal-bearing volumes, where the segmentation labels the existing implant as bone, treating the implant as occupied space (local half-maximum threshold per object) lowered 10 mm viability from 53.8% to 40.4% in 52 patients with the default crop, and a fixed 2500 HU threshold, which thins the screws and therefore bounds the occupied corridor from above, gave the same proportion; these volumes are reported as a separate contrast and are not pooled with the corridor cohort.

Local cohort audit. Training eligibility requires manual review of three-dimensional views and multiplanar CT slices confirming absence of metal and foreign objects. HU screening only prioritizes this review, including cases from dataset6. Confirmed metal cases are candidates for testing, with pelvic osteosynthesis distinguished from other implants and external objects. Exact volume duplicates and repeated patients must be resolved before splitting. Final cohort sizes remain pending this audit.

Why downstream evaluation is excluded. The audit itself supplies the reason. Only a minority of the metal-bearing volumes carry verified bone annotation, and only a fraction of the published collection is available locally, so a Dice or HD95 result would rest on a cohort too small and too partially annotated to support a robustness claim. The metal-bearing subset used here is nonetheless the de facto clinical evaluation set for pelvic metal-artifact work, having been used as such in recent published work (Wang et al., 2025). Where a clean reference image does not exist for these volumes, appearance realism is assessed against real artifact observations rather than against an unavailable ground truth.

7. References

- E. Abadi, B. Harrawood, S. Sharma, A. Kapadia, W. P. Segars, and E. Samei. DukeSim: A Realistic, Rapid, and Scanner-Specific Simulation Framework in Computed Tomography. *IEEE Transactions on Medical Imaging*, 38(6):1457–1465, 2019. doi: 10.1109/TMI.2018.2886530.
- C. Arand, D. Wagner, R. G. Richards, H. Noser, L. Kamer, T. Sawaguchi, and P. M. Rommens. 3D Statistical Model of the Pelvic Ring — A CT-Based Statistical Evaluation of Anatomical Variation. *Journal of Anatomy*, 234(3):376–383, 2019. doi: 10.1111/joa.12928.
- Q. Chen, X. Chen, H. Song, Z. Xiong, A. Yuille, C. Wei, and Z. Zhou. Towards Generalizable Tumor Synthesis. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11147–11158, 2024. doi: 10.1109/ICCV52733.2024.01060.
- Q. Chen, S. Ding, Y. Gu, N. Liu, J. Bian, A. Yuille, Z. Zhou, and J. Fu. Foundation VAEs for 3D CT Reconstruction, Augmentation, and Generation, 2026.
- B. De Man, J. Nuyts, P. Dupont, G. Marchal, and P. Suetens. Metal Streak Artifacts in X-Ray Computed Tomography: A Simulation Study. *IEEE Transactions on Nuclear Science*, 46(3):691–696, 1999. doi: 10.1109/23.775600.
- B. De Man, S. Basu, N. Chandra, B. Dunham, P. Edic, M. Iatrou, S. McOlash, P. Sainath, C. Shaughnessy, B. Tower, and E. Williams. CatSim: A New Computer Assisted Tomography Simulation Environment. In J. Hsieh and M. J. Flynn, editors, *Medical Imaging 2007: Physics of Medical Imaging*, volume 6510, page 65102G. SPIE, 2007. doi: 10.1117/12.710713.
- Z. Dorjsenbe, H.-K. Pao, and F. Xiao. Polyp-DDPM: Diffusion-Based Semantic Polyp Synthesis for Enhanced Segmentation. In *2024 46th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 1–7, 2024. doi: 10.1109/EMBC53108.2024.10782077.
- N. A. Ebraheim, R. J. Coombs, M. J. Hoefflinger, and W. T. Jackson. A Pitfall of Radiologic Evaluation of Sacroiliac Joint Screw Positioning. *Orthopedics*, 16(5):616–618, 1993. doi: 10.3928/0147-7447-19930501-16.
- N. A. Ebraheim, R. Xu, A. Biyani, and M. C. Nadaud. Morphologic Considerations of the First Sacral Pedicle for Iliosacral Screw Placement. *Spine*, 22(8):841–846, 1997. doi: 10.1097/00007632-199704150-00002.
- Y. Fan, J. Pack, and B. De Man. A Virtual Imaging Trial Framework to Study Cardiac CT Blooming Artifacts. In *7th International Conference on Image Formation in X-Ray Computed Tomography*, volume 12304 of *Society of Photo-Optical Instrumentation Engineers (SPIE) Conference Series*, page 1230415, 2022. doi: 10.1117/12.2646407.
- M. J. Gardner, S. Morshed, S. E. Nork, W. M. Ricci, and M. L. Chip Routt, Jr. Quantification of the Upper and Second Sacral Segment Safe Zones in Normal and Dysmorphic Sacra. *Journal of Orthopaedic Trauma*, 24(10):622–629, 2010. doi: 10.1097/BOT.0b013e3181cf0404.
- S. D. Gertzbein and S. E. Robbins. Accuracy of Pedicular Screw Placement In Vivo. *Spine*, 15(1):11–14, 1990. doi: 10.1097/00007632-199001000-00004.
- H. Gottschling, M. Schröder, N. Reimers, F. Fischer, A. Homeier, and R. Burgkart. A System for Performing Automated Measurements on Large Bone Databases. In O. Dössel and W. C. Schlegel, editors, *World Congress on Medical Physics and Biomedical Engineering, September 7 - 12, 2009, Munich, Germany*, pages 910–913, Berlin, Heidelberg, 2010. Springer Berlin Heidelberg. ISBN 978-3-642-03882-2.
- F. Gras, H. Gottschling, M. Schröder, I. Marintschev, G. O. Hofmann, and R. Burgkart. Transsacral Osseous Corridor Anatomy Is More Amenable To Screw Insertion In Males: A Biomorphometric Analysis of 280 Pelvis. *Clinical Orthopaedics and Related Research*, 474(10):2304–2311, 2016. doi: 10.1007/s11999-016-4954-5.
- E. Haneda, N. Peters, J. Zhang, G. Karageorgos, W. Xia, H. Paganetti, G. Wang, Y. Guo, J. Ma, H. S. Park, K. Jeon, F. Fan, M. Thies, and B. De Man. AAPM CT Metal Artifact Reduction Grand Challenge. *Medical Physics*, 52(10):e70050, 2025. doi: 10.1002/mp.70050.
- E. A. Hasenboehler, P. F. Stahel, A. Williams, W. R. Smith, J. T. Newman, D. L. Symonds, and S. J. Morgan. Prevalence of Sacral Dysmorphism in a Prospective Trauma Population: Implications for a “Safe” Surgical Corridor for Sacro-Iliac Screw Placement. *Patient Safety in Surgery*, 5(1):8, 2011. doi: 10.1186/1754-9493-5-8.
- A. Herman, E. Keener, C. Dubose, and J. A. Lowe. Simple Mathematical Model of Sacroiliac Screws Safe-Zone-Easy to Implement by Pelvic Inlet and Outlet Views. *Journal of Orthopaedic Research*, 35(7):1478–1484, 2017. doi: 10.1002/jor.23396.
- A. F. Hinsche, P. V. Giannoudis, and R. M. Smith. Fluoroscopy-Based Multiplanar Image Guidance for Insertion of Sacroiliac Screws. *Clinical Orthopaedics and Related Research*, (395):135–144, 2002. doi: 10.1097/00003086-200202000-00014.
- Q. Hu, Y. Chen, J. Xiao, S. Sun, J. Chen, A. Yuille, and Z. Zhou. Label-Free Liver Tumor Segmentation. In *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7422–7432, 2023. doi: 10.1109/CVPR52729.2023.00717.
- F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein. nnU-Net: A Self-Configuring Method for Deep Learning-Based Biomedical Image Segmentation. *Nature Methods*, 18(2):203–211, 2021. doi: 10.1038/s41592-020-01008-z.
- A. J. Jacob, P. Sharma, and D. Rueckert. LGESynthNet: Controlled Scar Synthesis for Improved Scar Segmentation in Cardiac LGE-MRI Imaging. In O. Camara, E. Puyol Antón, M. Sermesant, C. Mauge, M. Varela, Y. Ma, R. Paulsen, C. Wang, Q. Tao, and A. Young, editors, *Statistical Atlases and Computational Models of the Heart. Regular and CMR-Recon Challenge Papers*, pages 34–44, Cham, 2026. Springer Nature Switzerland. ISBN 978-3-032-17734-6. doi: 10.1007/978-3-032-17734-6_4.
- S. P. Kaiser, M. J. Gardner, J. Liu, M. L. C. Routt, Jr., and S. Morshed. Anatomic Determinants of Sacral Dysmorphism and Implications for Safe Iliosacral Screw Placement. *The Journal of Bone and Joint Surgery. American Volume*, 96(14):e120, 2014. doi: 10.2106/JBJS.M.00895.
- G. M. Karageorgos, J. Zhang, N. Peters, W. Xia, C. Niu, H. Paganetti, G. Wang, and B. De Man. A Denoising Diffusion Probabilistic Model for Metal Artifact Reduction in CT. *IEEE Transactions on Medical Imaging*, 43(10):3521–3532, 2024. doi: 10.1109/TMI.2024.3416398.
- A. Kazerouni, E. K. Aghdam, M. Heidari, R. Azad, M. Fayyaz, I. Hachililoglu, and D. Merhof. Diffusion Models in Medical Imaging: A Comprehensive Survey. *Medical Image Analysis*, 88:102846, 2023. doi: 10.1016/j.media.2023.102846.
- J. F. Keating, J. Werier, P. Blachut, H. Broekhuysen, R. N. Meek, and P. J. O’Brien. Early Fixation of the Vertically Unstable Pelvis: The Role of Iliosacral Screw Fixation of the Posterior Lesion. *Journal of Orthopaedic Trauma*, 13(2):107–113, 1999. doi: 10.1097/00005131-199902000-00007.
- J. J. Lee, S. L. Rosenbaum, A. Martusiewicz, S. A. Holcombe, S. C. Wang, and J. A. Goulet. Transsacral Screw Safe Zone Size by Sacral Segmentation Variations. *Journal of Orthopaedic Research*, 33(2):277–282, 2015. doi: 10.1002/jor.22739.
- Z. Li, Q. Gao, Y. Wu, C. Niu, J. Zhang, M. Wang, G. Wang, and H. Shan. Quad-Net: Quad-Domain Network for CT Metal Artifact Reduction. *IEEE Transactions on Medical Imaging*, 43(5):1866–1879, 2024. doi: 10.1109/TMI.2024.3351722.
- W.-A. Lin, H. Liao, C. Peng, X. Sun, J. Zhang, J. Luo, R. Chellappa, and S. K. Zhou. DuDoNet: Dual Domain Network for CT Metal Artifact Reduction. In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10504–10513, 2019. doi: 10.1109/CVPR.2019.01076.
- J. Liu, H. Li, B. Zeng, H. Wang, R. Kikinis, L. Joskowicz, and X. Chen. An End-to-End Geometry-Based Pipeline for Automatic Preoperative Surgical Planning of Pelvic Fracture Reduction and Fixation. *IEEE Transactions on Medical Imaging*, 44(1):79–91, 2025. doi: 10.1109/TMI.2024.3429403.
- P. Liu, H. Han, Y. Du, H. Zhu, Y. Li, F. Gu, H. Xiao, J. Li, C. Zhao, L. Xiao, X. Wu, and S. K. Zhou. Deep Learning to Segment Pelvic Bones: Large-Scale CT Datasets and Baseline Models. *International Journal of Computer Assisted Radiology and Surgery*, 16(5):749–756, 2021. doi: 10.1007/s11548-021-02363-8.
- R. Macháček, L. Mozaffari, Z. Sepasdar, S. Parasa, P. Halvorsen, M. A. Riegler, and V. Thambawita. Mask-Conditioned Latent Diffusion for Generating Gastrointestinal Polyp Images. In *Proceedings of the 4th ACM Workshop on Intelligent Cross-Data Analysis and Retrieval, ICDAR ’23*, pages 1–9, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9789400701863. doi: 10.1145/3592571.3592978.
- D. A. McLaren, G. A. Busel, H. R. Parikh, A. Only, J. Patterson, B. T. Gaston, R. McLemore, and B. Cunningham. Corridor-Diameter-Dependent Angular Tolerance for Safe Transiliosacral Screw Placement: An Anatomic Study of 433 Pelvis. *European Journal of Orthopaedic Surgery & Traumatology*, 31(7):1485–1492, 2021. doi: 10.1007/s00590-021-02913-5.
- T. Mendel, H. Noser, D. Wohlrab, K. Stock, and F. Radetzki. The Lateral Sacral Triangle—A Decision Support for Secure Transverse Sacroiliac Screw Insertion. *Injury*, 42(10):1164–1170, 2011. doi: 10.1016/j.injury.2010.03.016.
- S. K. Mirza, G. C. Wiggins, C. Kuntz, J. E. York, C. Bellabarba, M. A. Knonodi, J. R. Chapman, and C. I. Shaffrey. Accuracy of Thoracic Vertebral Body Screw Placement Using Standard Fluoroscopy, Fluoroscopic Image Guidance, and Computed Tomographic Image Guidance: A Cadaver Study. *Spine*, 28(4):402–413, 2003. doi: 10.1097/01.BRS.0000048461.51308.CD.
- B. R. Moed and B. L. Geer. S2 Iliosacral Screw Fixation for Disruptions of the Posterior Pelvic Ring: A Report of 49 Cases. *Journal of Orthopaedic Trauma*, 20(6):378–383, 2006. doi: 10.1097/00005131-200607000-00002.
- N. Peters, E. Haneda, J. Zhang, G. Karageorgos, W. Xia, J. Verburg, G. Wang, H. Paganetti, and B. De Man. A Hybrid Training Database and Evaluation Benchmark for Assessing Metal Artifact Reduction Methods for X-ray CT Imaging. *Medical Physics*, 52(10):e70020, 2025. doi: 10.1002/mp.70020.
- N. Ramadanov and S. Zabler. Ramadanov-Zabler Safe Zone for Sacroiliac Screw Placement: A CT-Based Computational Pilot Study. *Journal of Clinical Medicine*, 14(10):3567, 2025. doi: 10.3390/jcm14103567.
- F. Ramzan, Y. Kiberu, N. Jathanna, S. Jamil-Copley, R. H. Clayton, and C. Chen. CLAIM: Clinically-Guided LGE Augmentation for Realistic and Diverse Myocardial Scar Synthesis and Segmentation. In D. Cafolla, T. Rittman, and H. Ni, editors, *Artificial Intelligence in Healthcare*, pages 279–292, Cham, 2026. Springer Nature Switzerland. ISBN 978-3-032-00652-3. doi: 10.1007/978-3-032-00652-3_20.
- L. Ren, A. Ferrero, and C. Favazza. Framework of Metal Insertion in the Projection Domain for Image Quality Optimization in Interventional Computed Tomography. *Journal of Medical Imaging*, 9(3):035001, 2022. doi: 10.1117/1.JMI.9.3.035001.
- R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-Resolution Image Synthesis with Latent Diffusion Models. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10674–10685, 2022. doi: 10.1109/CVPR52688.2022.01042.
- M. L. C. Routt, P. T. Simonian, and J. Inaba. Iliosacral Screw Complications. *Operative Techniques in Orthopaedics*, 7(3):206–220, 1997. doi: 10.1016/S1048-6666(97)80060-2. Pelvic Trauma: An update.
- M. Selles, J. A. C. van Osch, M. Maas, M. F. Boomsma, and R. H. H. Wellenb. Advances in Metal Artifact Reduction in CT Images: A Review of Traditional and Novel Metal Artifact Reduction Techniques. *European Journal of Radiology*, 170:111276, 2024. doi: 10.1016/j.ejrad.2023.111276.
- K. Singhrao, M. Zubair, T. Nano, J. E. Scholey, and M. Descovich. End-to-End Validation of Fiducial Tracking Accuracy in Robotic Radiosurgery Using MRI-Only Simulation Imaging. *Medical Physics*, 51(1):31–41, 2024. doi: 10.1002/mp.16857.
- H. E. Smith, P. S. Yuan, R. Sasso, S. Papadopoulos, and A. R. Vaccaro. An Evaluation of Image-Guided Technologies in the Placement of Percutaneous Iliosacral Screws. *Spine*, 31(2):234–238, 2006. doi: 10.1097/01.brs.00000194788.45002.1b.
- N. C. Tejwani, D. Raskolnikov, T. McLaurin, and R. Takemoto. The Role of Computed Tomography for Postoperative Evaluation of Percutaneous Sacroiliac Screw Fixation and Description of a “Safe Zone”. *American Journal of Orthopaedics*, 43(11):513–516, 2014.
- D. Templeman, A. Schmidt, J. Freese, and I. Weisman. Proximity of Iliosacral Screws to Neurovascular Structures After Internal Fixation. *Clinical Orthopaedics and Related Research*, (329):194–198, 1996. doi: 10.1097/00003086-199608000-00023.
- H. J. P. van Bosse, D. Lee, E. R. Henderson, D. A. Sala, and D. S. Feldman. Pelvic Positioning Creates Error in CT Acetabular Measurements. *Clinical Orthopaedics and Related Research*, 469(6):1683–1691, 2011. doi: 10.1007/s11999-011-1827-9.
- E. W. van den Bosch, C. M. A. van Zwiene, and A. B. van Vugt. Fluoroscopic Positioning of Sacroiliac Screws in 88 Patients. *The Journal of Trauma*, 53(1):44–48, 2002. doi: 10.1097/00005373-200207000-00009.
- D. Wagner, L. Kamer, T. Sawaguchi, R. Geoff Richards, H. Noser, M. Uesugi, C. Ossendorf, and P. M. Rommens. Critical Dimensions of Trans-Sacral Corridors Assessed by 3D CT Models: Relevance for Implant Positioning in Fractures of the Sacrum. *Journal of Orthopaedic Research*, 35(11):2577–2584, 2017. doi: 10.1002/jor.23554.
- H. Wang, Y. Wu, Y. Wang, D. Wei, X. Wu, J. Ma, and Y. Zheng. Adaptive Weighting Based Metal Artifact Reduction in CT Images. *IEEE Transactions on Medical Imaging*, 44(6):2408–2423, 2025. doi: 10.1109/TMI.2025.3534316.
- Z. Wang, C. Vandersteen, T. Demarcy, D. Gnansia, C. Raffaelli, N. Guevara, and H. Delingette. Deep Learning Based Metal Artifacts Reduction in Post-operative Cochlear Implant CT Imaging. In D. Shen, T. Liu, T. M. Peters, L. H. Staib, C. Essert, S. Zhou, P.-T. Yap, and A. Khan, editors, *Medical Image Computing and Computer Assisted Intervention — MICCAI 2019*, pages 121–129, Cham, 2019. Springer International Publishing. ISBN 978-3-030-32226-7. doi: 10.1007/978-3-030-32226-7_14.
- J. Wasserthal, H.-C. Breit, M. T. Meyer, M. Pradella, D. Hinck, A. W. Sauter, T. Heye, D. T. Boll, J. Cyriac, S. Yang, M. Bach, and M. Segeroth. TotalSegmentor: Robust Segmentation of 104 Anatomic Structures in CT Images. *Artificial Intelligence*, 5(5):e230024, 2023. doi: 10.1148/ryai.230024.
- M. Wu, P. FitzGerald, J. Zhang, W. P. Segars, H. Yu, Y. Xu, and B. De Man. XCIST — An Open Access X-ray/CT Simulation Toolkit. *Physics in Medicine & Biology*, 67(19):194002, 2022. doi: 10.1088/1361-6560/ac9174.
- K. Xie, L. Gao, Y. Zhang, H. Zhang, J. Sun, T. Lin, J. Sui, and X. Ni. Metal Implant Segmentation in CT Images Based on Diffusion Model. *BMC Medical Imaging*, 24(1):204, 2024. doi: 10.1186/s12880-024-01379-1.
- L. Yu, Z. Zhang, X. Li, and L. Xing. Deep Sinogram Completion With Image Prior for Metal Artifact Reduction in CT Images. *IEEE Transactions on Medical Imaging*, 40(1):228–238, 2021. doi: 10.1109/TMI.2020.3025064.
- S. Yun, S. Hyun, D.-I. Choi, and S. Cho. A Strategy for Simulation-Driven CT Metal Artifact Reduction Toward Improving Network Generalizability. *Medical Physics*, 53(2):e70336, 2026. doi: 10.1002/mp.70336.
- L. Zhang, A. Rao, and M. Agrawala. Adding Conditional Control to Text-to-Image Diffusion Models. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3813–3824, 2023. doi: 10.1109/ICCV51070.2023.00355.
- Y. Zhang and H. Yu. Convolutional Neural Network Based Metal Artifact Reduction in X-Ray Computed Tomography. *IEEE Transactions on Medical Imaging*, 37(6):1370–1381, 2018. doi: 10.1109/TMI.2018.2823083.
- Y. Zhang, H. Luo, D. He, Z. Guan, J. Zhao, and Z. Yang. Enhancing Intraoperative Pedicle Screw Planning Accuracy via Diffusion-Based Synthetic CT. *Journal of X-Ray Science and Technology*, 34(4):525–540, 2026. doi: 10.1177/08953996261443500.
- Z. Zhang, L. Yao, B. Wang, D. Jha, G. Durak, E. Keles, A. Medetalibeyoglu, and U. Bagci. DiffBoost: Enhancing Medical Image Segmentation via Text-Guided Diffusion Model. *IEEE Transactions on Medical Imaging*, 44(9):3670–3682, 2025. doi: 10.1109/TMI.2024.3519307.
- Y. Zhao, J. Li, D. Wang, and W. Lian. Parameters of Lengthened Sacroiliac Screw Fixation: A Radiological Anatomical Study. *European Spine Journal*, 21(9):1807–1814, 2012. doi: 10.1007/s00586-012-2367-z.
- B. H. Ziran, A. D. Wasan, D. M. Marks, S. A. Olson, and M. W. Chapman. Fluoroscopic Imaging Guides of the Posterior Pelvis Pertaining to Iliosacral Screw Placement. *The Journal of Trauma*, 62(2):347–356, 2007. doi: 10.1097/01.ta.0000229789.18211.85.

J. Zwingmann, G. Konrad, E. Kotter, N. P. Südkamp, and M. Oberst. Computer-Navigated Iliosacral Screw Insertion Reduces Malposition Rate and Radiation Exposure. *Clinical Orthopaedics and Related Research*, 467(7):1833–1838, 2009. doi: 10.1007/s11999-008-0632-6.

J. Zwingmann, O. Hauschild, G. Bode, N. P. Südkamp, and H. Schmal. Malposition and Revision Rates of Different Imaging Modalities for Percutaneous Iliosacral Screw Fixation Following Pelvic Fractures: A Systematic Review and Meta-Analysis. *Archives of Orthopaedic and Trauma Surgery*, 133(9):1257–1265, 2013. doi: 10.1007/s00402-013-1788-4.